

Deep Learning Hardware Deployment Internship

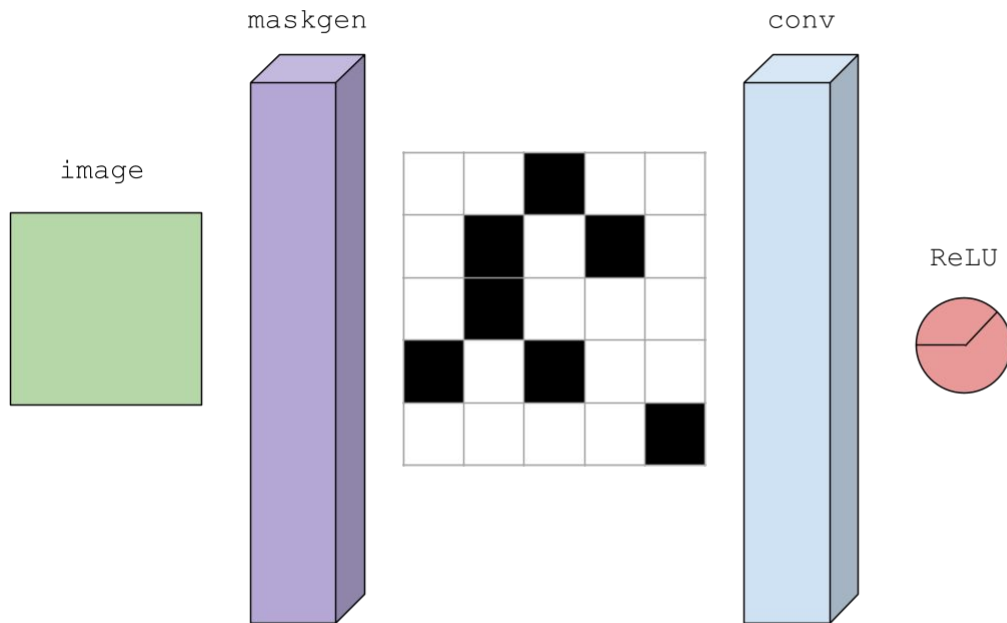
Zero Skipping Mask Generation

Young Jae Kim

2018-15653

Department of Electrical and Computer Engineering

Estimating Where the Zeros Will Be...



The CNN network

Layer	Filter	Input	Output	Mask
1	3x3	128x128x1	128x128x64	N/A
2	1x1	128x128x64	128x128x32	1x1x64x4, 1x1x4x32
3	3x3	128x128x32	128x128x32	1x1x32x4, 3x3x4x4, 1x1x4x32
4	3x3	128x128x32	128x128x32	1x1x32x4, 3x3x4x4, 1x1x4x32
5	3x3	128x128x32	128x128x32	1x1x32x4, 3x3x4x4, 1x1x4x32
6	3x3	128x128x32	128x128x32	1x1x32x4, 3x3x4x4, 1x1x4x32
7	1x1	128x128x32	128x128x64	1x1x32x4, 1x1x4x64
8	3x3	128x128x64	128x128x4	N/A

Is It Even Worth It?

Layer	Filter	Input	Output	Mask
3	3x3	128x128x32	128x128x32	1x1x32x4, 3x3x4x4, 1x1x4x32

To calculate the output feature map (OFM) requires

$$3 \times 3 \times 32 \times 32 \times 128 \times 128 = 150\text{M calculations}$$

To calculate the mask using our mask generation model architecture requires

$$128 \times 128 \times 32 \times 4 + 128 \times 128 \times 3 \times 3 \times 4 + 128 \times 128 \times 4 \times 32 = 4.8\text{M calculations}$$

Is It Even Worth It?

Layer	Filter	Input	Output	Mask
3	3x3	128x128x32	128x128x32	1x1x32x4, 3x3x4x4, 1x1x4x32

Calculation cost of calculating one 1x1x1 output in the OFM

$$3 \times 3 \times 32 \times 32 = 288$$

Number of zeros that need to be predicted in order for the calculation to “break even”

$$\frac{4784128(=\text{multiplications for mask calculation})}{288} = 16,611.6$$

$$\frac{16,611.6}{128 \times 128 \times 32(=\text{number pixels in OFM})} \approx 0.0317 = 3.17\%$$

At least 3.17% of the outputs should be 0 in order for the mask generation calculation to break even

Is It Even Worth It?

```
[ INFO ] layer 0 ( conv) (scale 00.16) (biases +-02.13) (bias_shift = 11 act_shift = 07) (nzs = 0.44)
[ INFO ] layer 1 ( conv) (scale 02.14) (biases +-02.13) (bias_shift = 15 act_shift = 07) (nzs = 0.36)
[ INFO ] layer 2 ( conv) (scale 00.16) (biases +-00.15) (bias_shift = 15 act_shift = 10) (nzs = 0.41)
[ INFO ] layer 3 ( conv) (scale 00.16) (biases +-00.15) (bias_shift = 14 act_shift = 09) (nzs = 0.60)
[ INFO ] layer 4 ( conv) (scale 00.16) (biases +-00.15) (bias_shift = 15 act_shift = 08) (nzs = 0.51)
[ INFO ] layer 5 ( conv) (scale 00.16) (biases +-00.15) (bias_shift = 16 act_shift = 07) (nzs = 0.57)
[ INFO ] layer 6 ( conv) (scale 00.16) (biases +-00.15) (bias_shift = 17 act_shift = 06) (nzs = 0.13)
[ INFO ] layer 7 ( conv) (scale 00.16) (biases +-00.15) (bias_shift = 18 act_shift = 07) (nzs = 0.78)
```

Is It Even Worth It?

[INFO]	layer	0	(conv)	(scale	00.16)	(biases +-02.13)	(bias_shift = 11	act_shift = 07)	(nzh = 0.44)
[INFO]	layer	1	(conv)	(scale	02.14)	(biases +-02.13)	(bias_shift = 15	act_shift = 07)	(nzh = 0.36)
[INFO]	layer	2	(conv)	(scale	00.16)	(biases +-00.15)	(bias_shift = 15	act_shift = 10)	(nzh = 0.41)
[INFO]	layer	3	(conv)	(scale	00.16)	(biases +-00.15)	(bias_shift = 14	act_shift = 09)	(nzh = 0.60)
[INFO]	layer	4	(conv)	(scale	00.16)	(biases +-00.15)	(bias_shift = 15	act_shift = 08)	(nzh = 0.51)
[INFO]	layer	5	(conv)	(scale	00.16)	(biases +-00.15)	(bias_shift = 16	act_shift = 07)	(nzh = 0.57)
[INFO]	layer	6	(conv)	(scale	00.16)	(biases +-00.15)	(bias_shift = 17	act_shift = 06)	(nzh = 0.13)
[INFO]	layer	7	(conv)	(scale	00.16)	(biases +-00.15)	(bias_shift = 18	act_shift = 07)	(nzh = 0.78)

All outputs from each layer have more than 3.17% of zero-ness

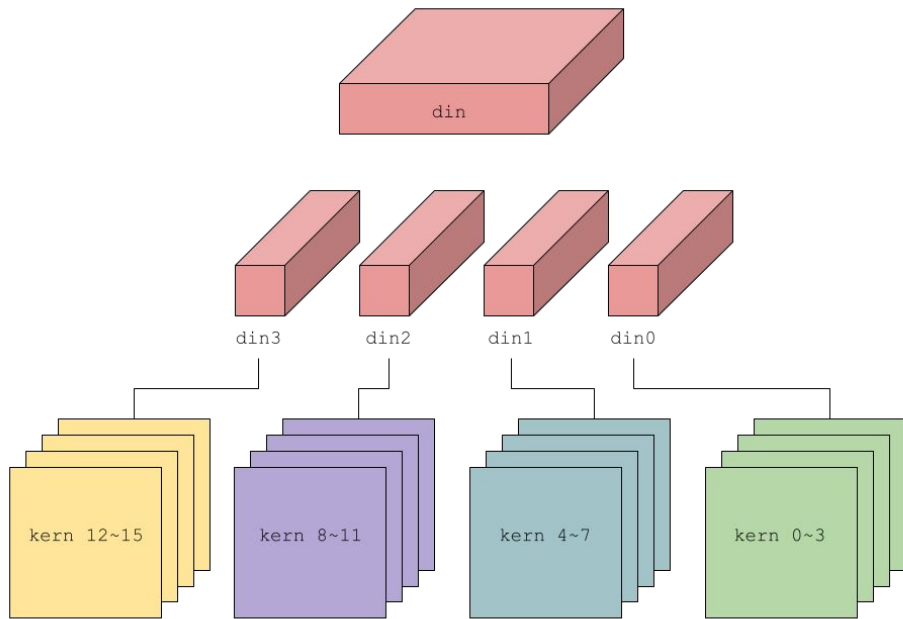
Yes, zero skipping via mask generation is worth it

Implementation Specifications/Constraints

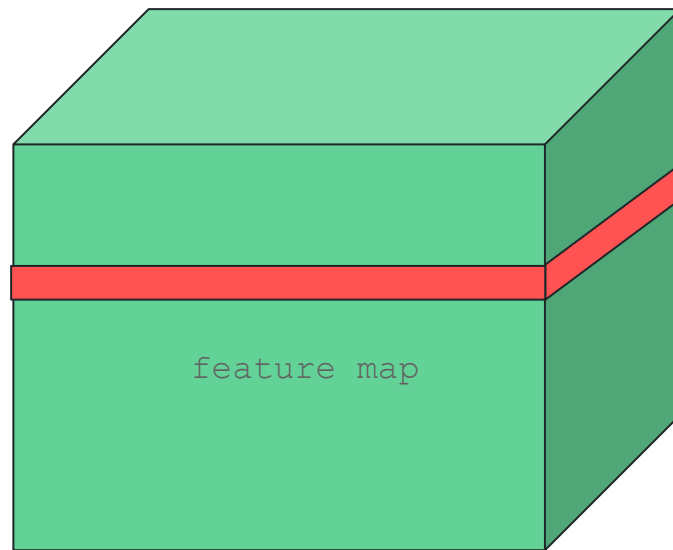
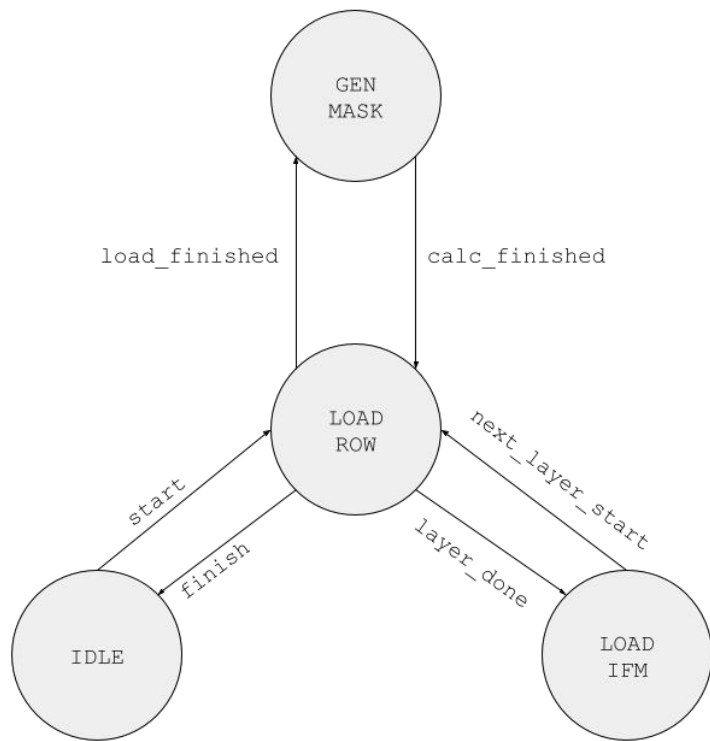
Mask_maker

Perform convolution with $1 \times 4 \times 4$ section (din) of feature map

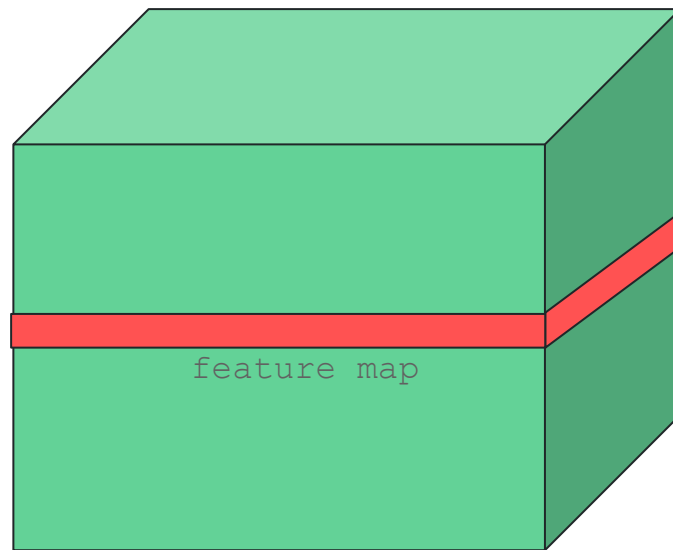
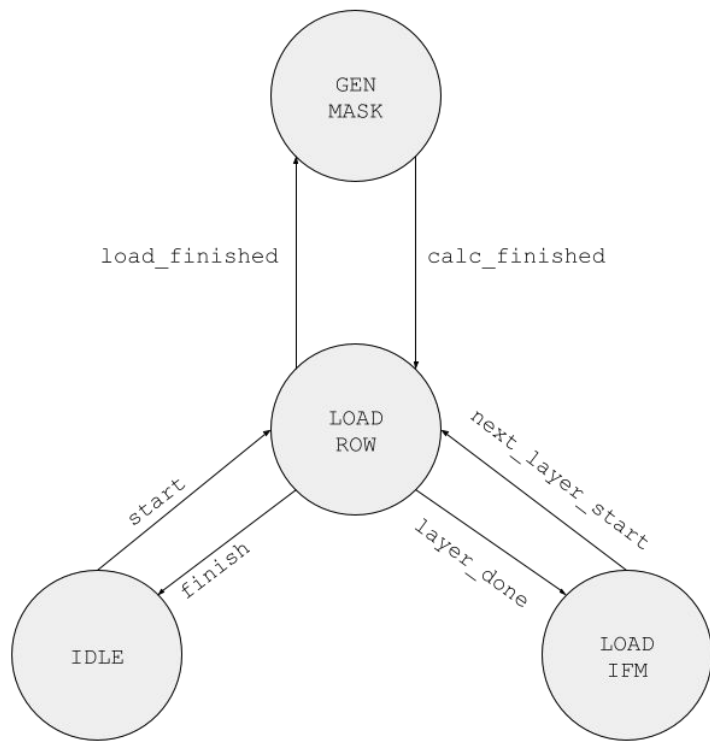
16 total kernels



Maskgen FSM



Maskgen FSM



Layer 1 Convolution Operations

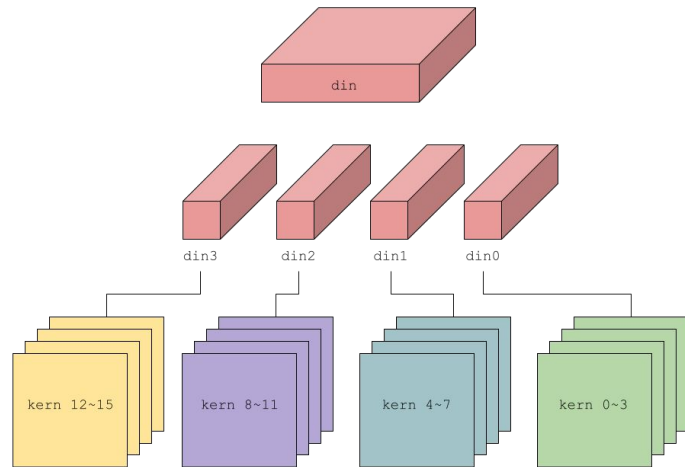
input : $128 \times 128 \times 32$

output : $128 \times 128 \times 4$

weight : $1 \times 1 \times 32 \times 4$

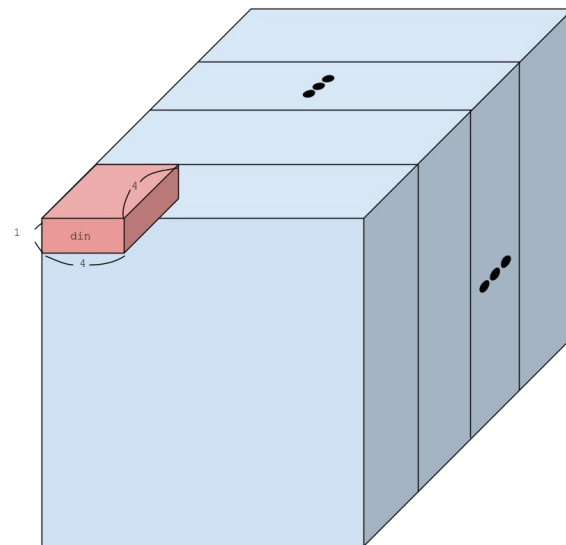
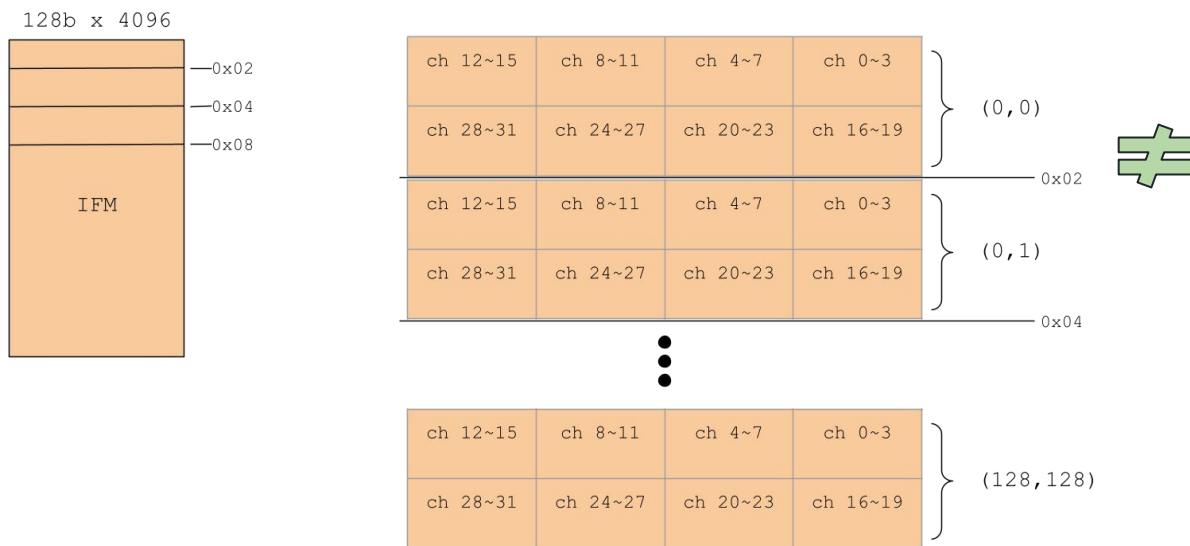
din only takes in 32 bits, so only 4 pixels ($4 \times 8 \text{ bit} = 32 \text{ bit}$) channel-wise enter at a time

Accumulate occurs after 8 “batches” of 4 pixels



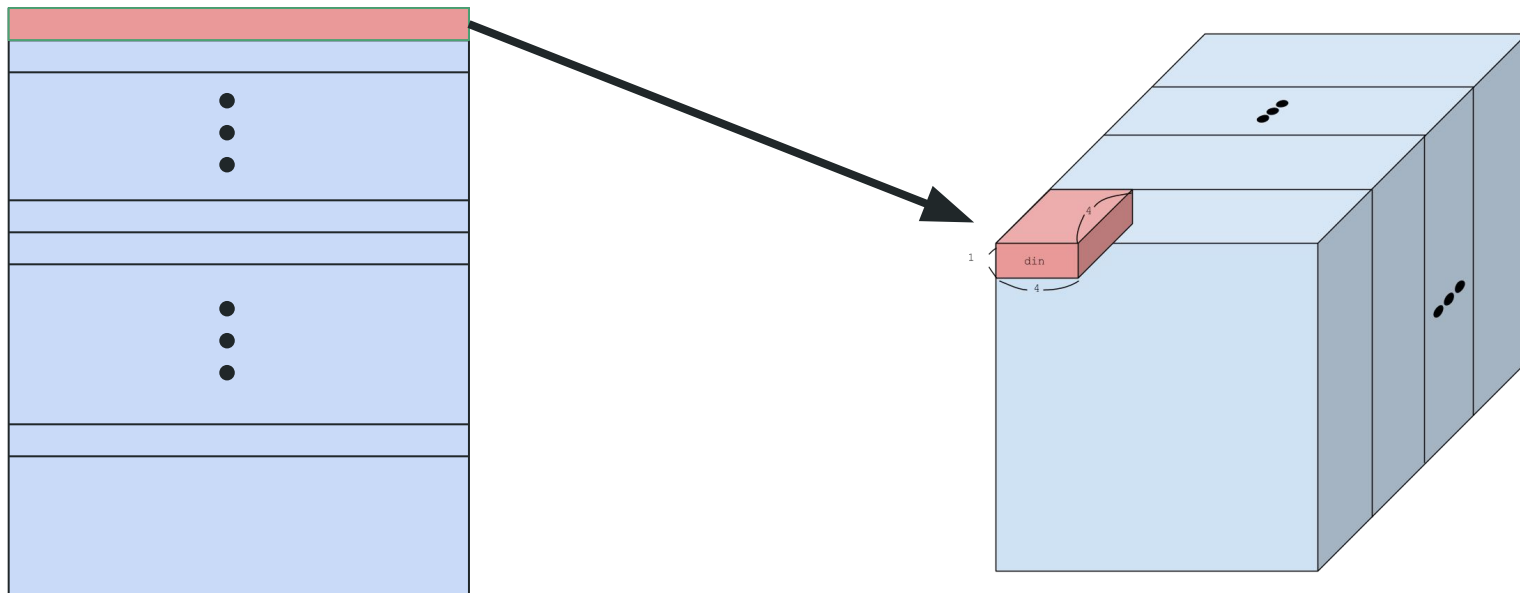
Inputs

Input feature map from SW code (convout)



Inputs

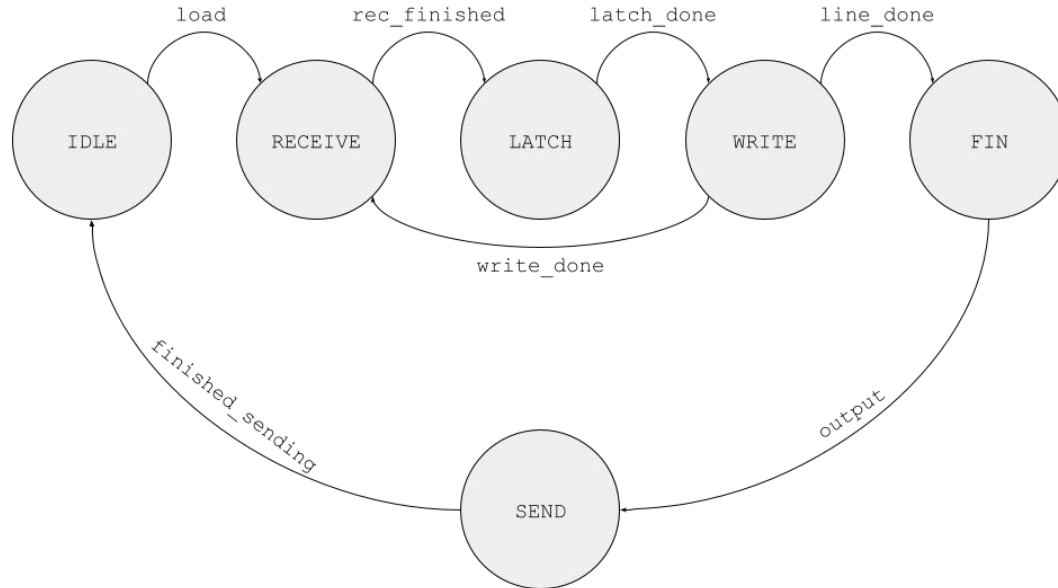
Input feature map from SW code (convout)



Layer 1 Line Buffer

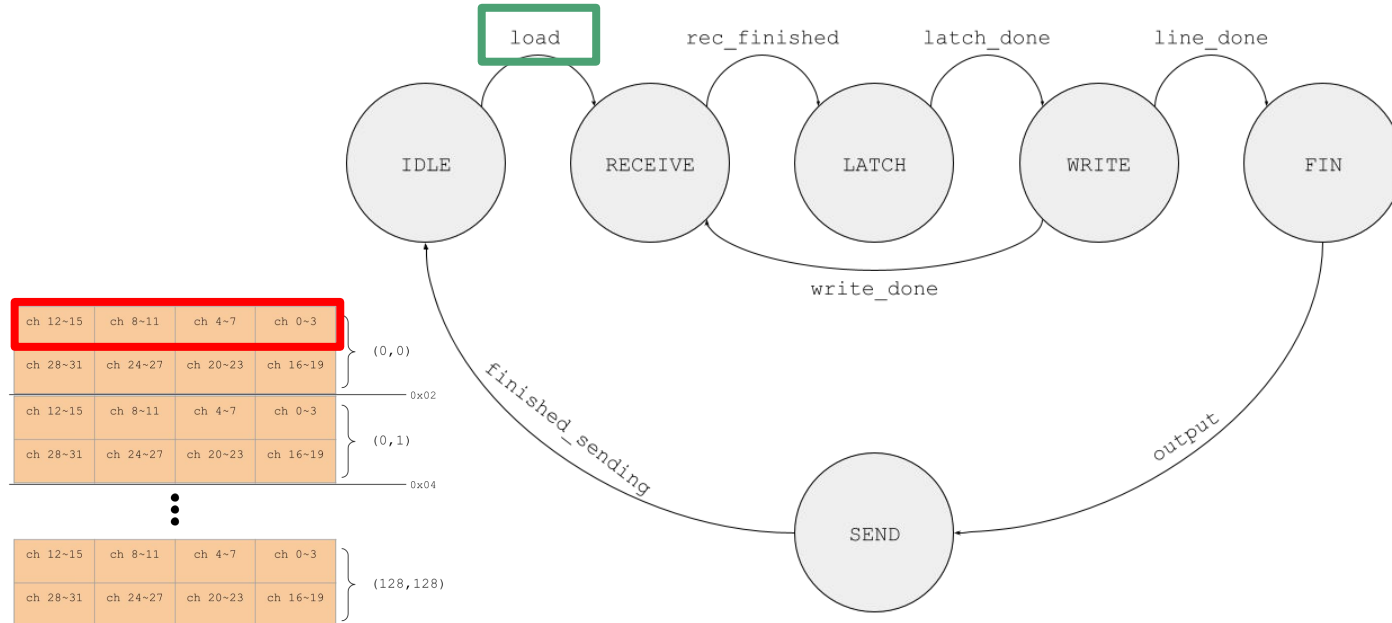
Holds 1 row worth of data (128x1x4 bytes)

Internally restructures data so 1 address holds 128b



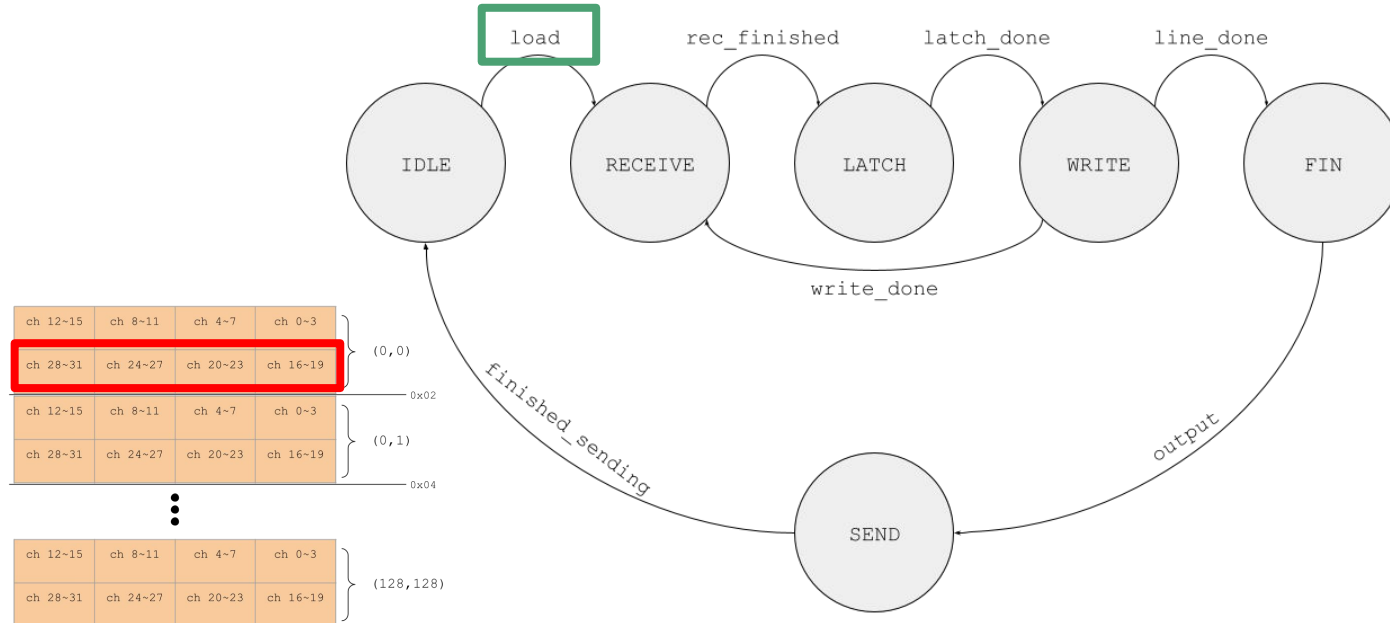
Layer 1 Line Buffer

In RECEIVE state, line buffer receives data(128b) from un-organized IFM.
Inputs stored in temporary buffer.



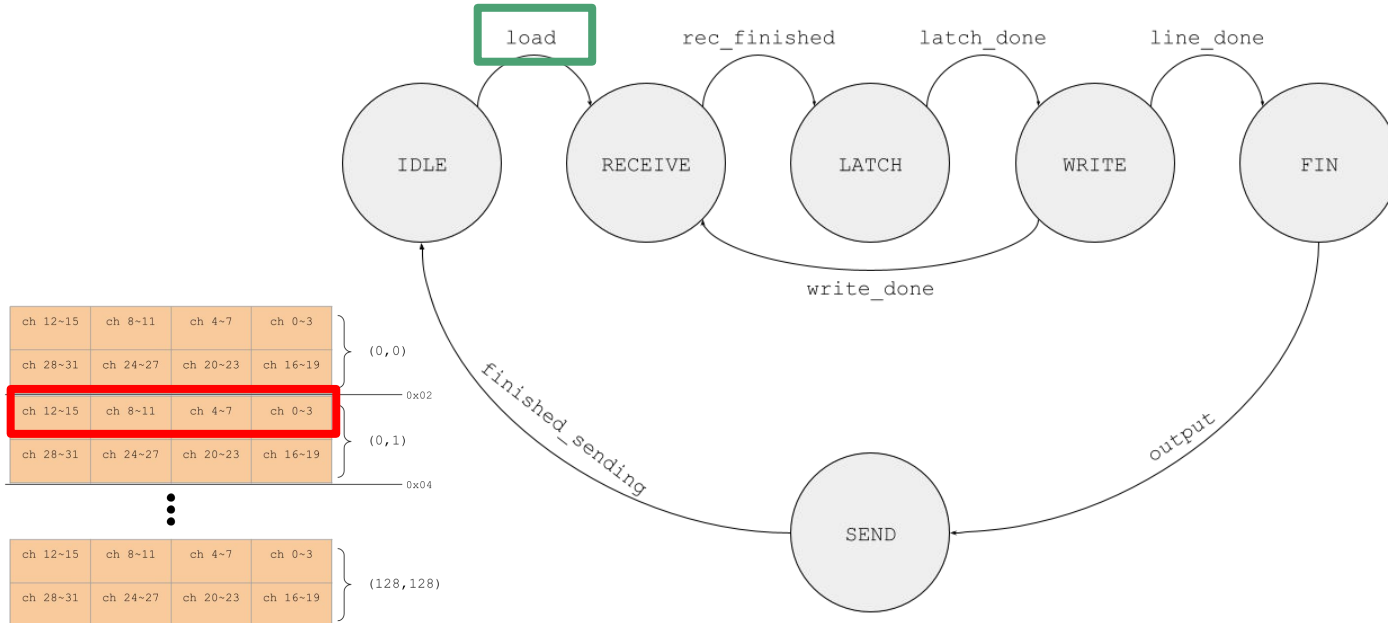
Layer 1 Line Buffer

In RECEIVE state, line buffer receives data(128b) from un-organized IFM.
Inputs stored in temporary buffer.



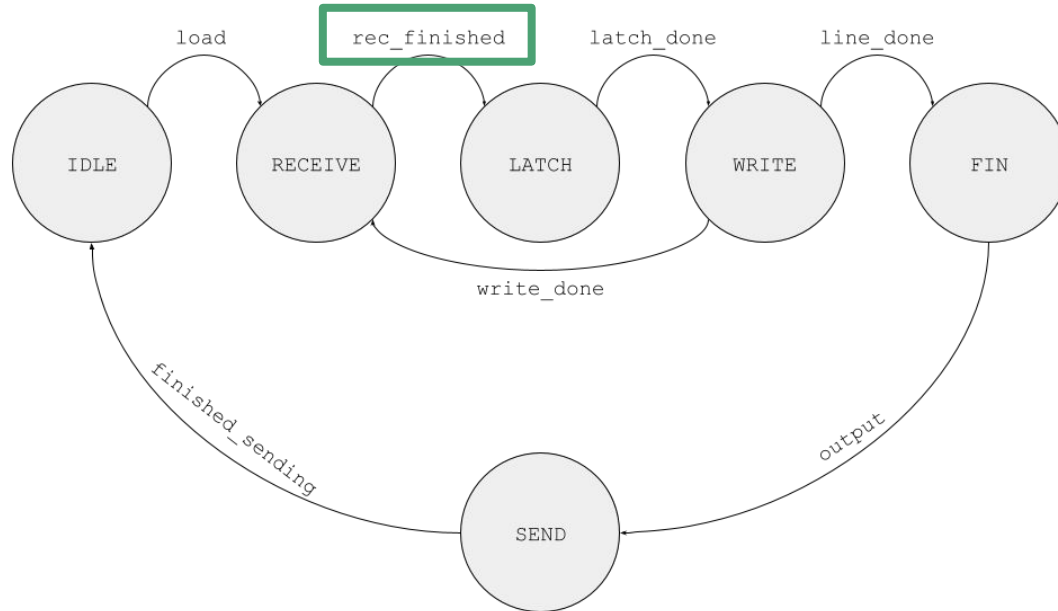
Layer 1 Line Buffer

In RECEIVE state, line buffer receives data(128b) from un-organized IFM.
Inputs stored in temporary buffer.



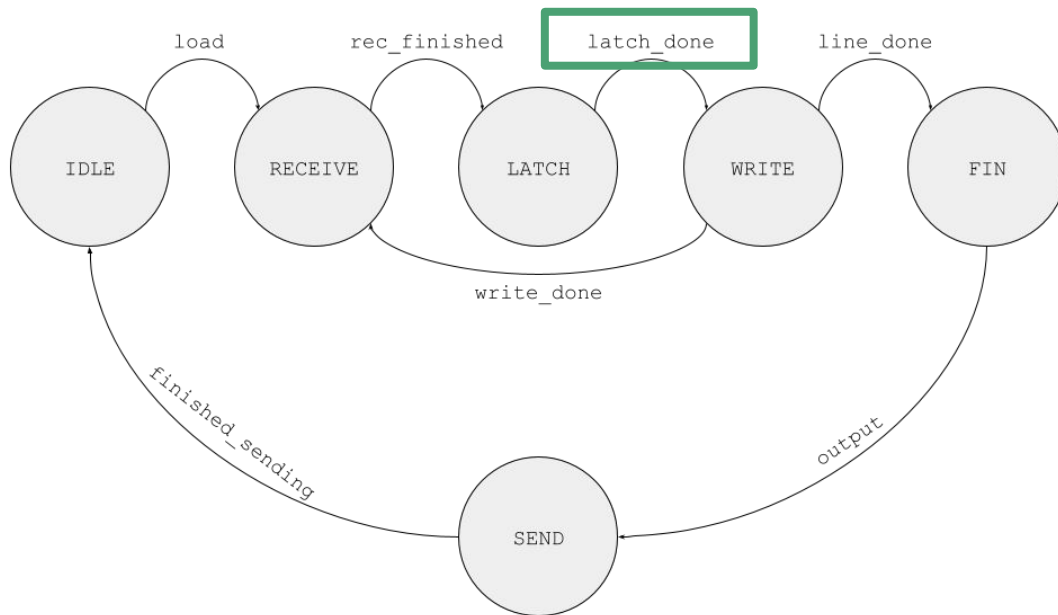
Layer 1 Line Buffer

When 128b from the IFM received 8 times, go to LATCH state



Layer 1 Line Buffer

In LATCH state, reorganize the received weights so that 1 address holds data for one 1x4x4 convolution.



Layer 1 Line Buffer

After weights reorganized, move onto WRITE state

```
if(cstate == ST_LATCH)begin
  if(receiving_line_cnt_d2 == 4'd8)begin
    tmp_write[0] <= {tmp_inputs[0][0*32+:32], tmp_inputs[2][0*32+:32], tmp_inputs[4][0*32+:32], tmp_inputs[6][0*32+:32]};
    tmp_write[1] <= {tmp_inputs[0][1*32+:32], tmp_inputs[2][1*32+:32], tmp_inputs[4][1*32+:32], tmp_inputs[6][1*32+:32]};
    tmp_write[2] <= {tmp_inputs[0][2*32+:32], tmp_inputs[2][2*32+:32], tmp_inputs[4][2*32+:32], tmp_inputs[6][2*32+:32]};
    tmp_write[3] <= {tmp_inputs[0][3*32+:32], tmp_inputs[2][3*32+:32], tmp_inputs[4][3*32+:32], tmp_inputs[6][3*32+:32]};

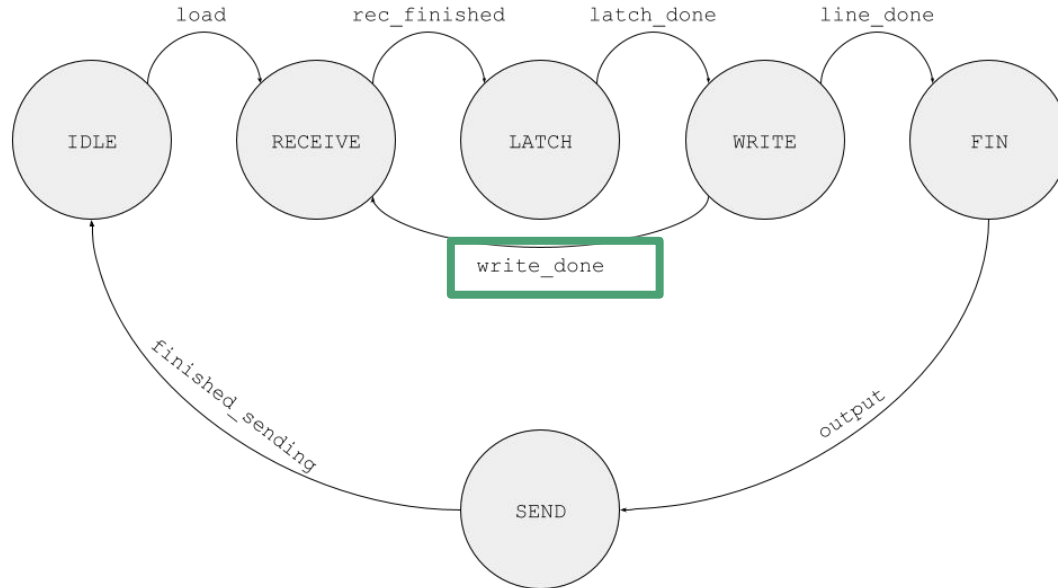
    tmp_write[4] <= {tmp_inputs[1][0*32+:32], tmp_inputs[3][0*32+:32], tmp_inputs[5][0*32+:32], tmp_inputs[7][0*32+:32]};
    tmp_write[5] <= {tmp_inputs[1][1*32+:32], tmp_inputs[3][1*32+:32], tmp_inputs[5][1*32+:32], tmp_inputs[7][1*32+:32]};
    tmp_write[6] <= {tmp_inputs[1][2*32+:32], tmp_inputs[3][2*32+:32], tmp_inputs[5][2*32+:32], tmp_inputs[7][2*32+:32]};
    tmp_write[7] <= {tmp_inputs[1][3*32+:32], tmp_inputs[3][3*32+:32], tmp_inputs[5][3*32+:32], tmp_inputs[7][3*32+:32]};

    latch_done <= 1;
  end
end
```

Layer 1 Line Buffer

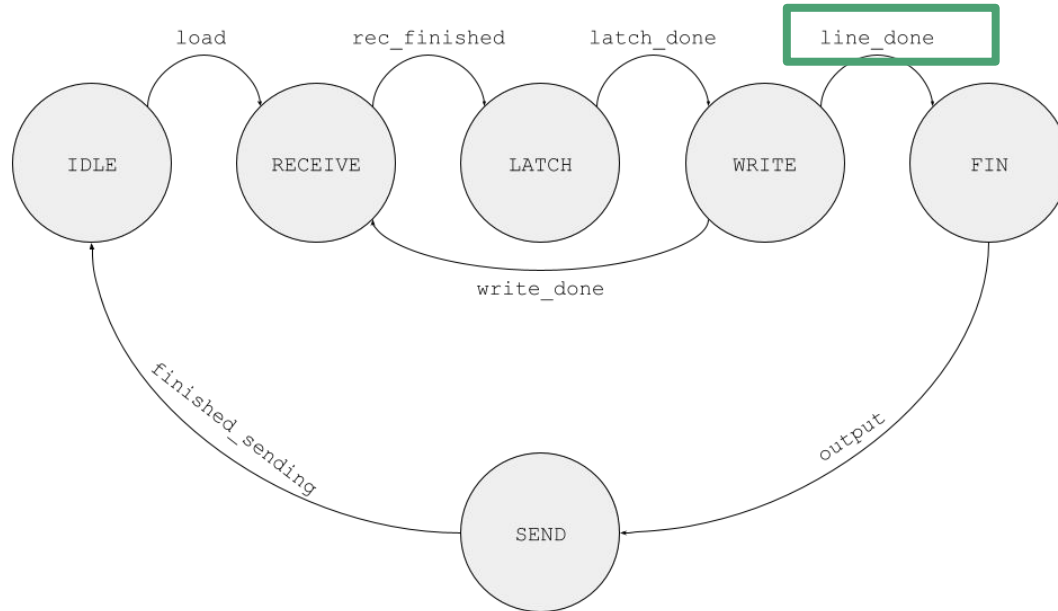
In WRITE state, write the 8 lines of restructured data.

If the entire row is not yet complete, go to RECEIVE state and repeat



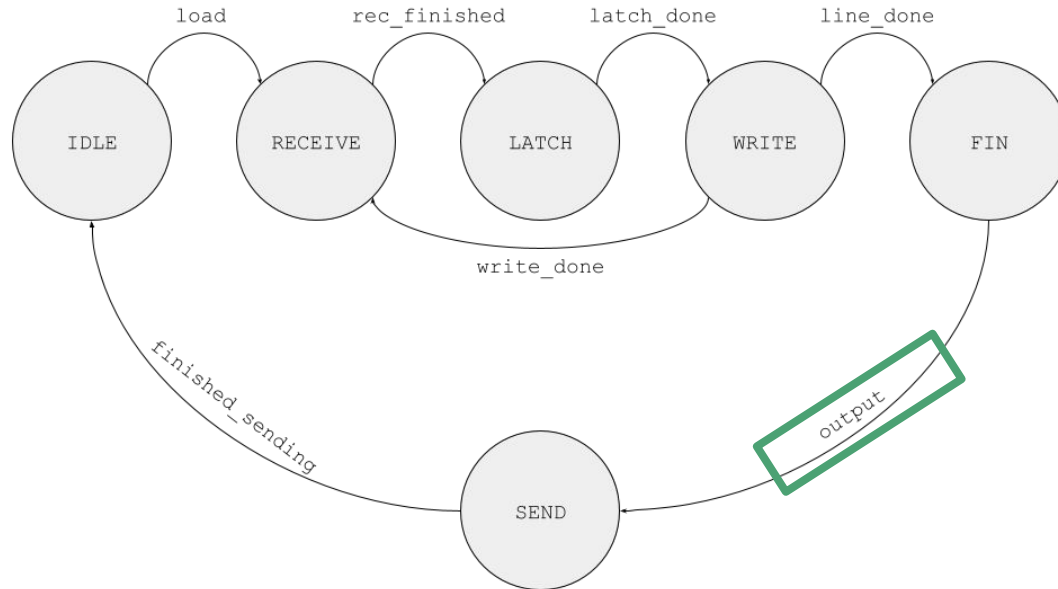
Layer 1 Line Buffer

If entire row has been received and written, go to FIN state (finished receiving)



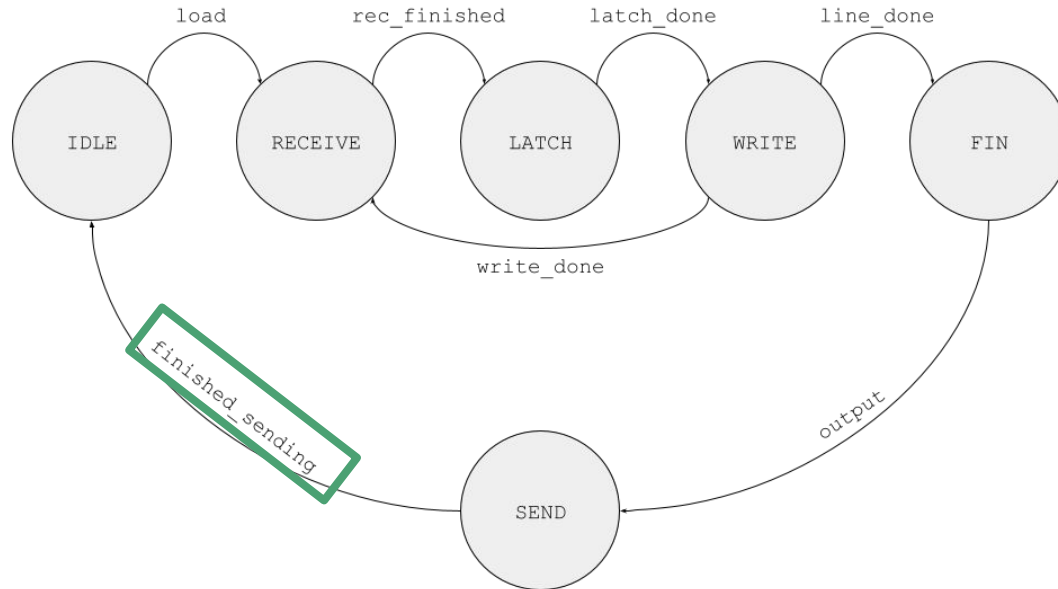
Layer 1 Line Buffer

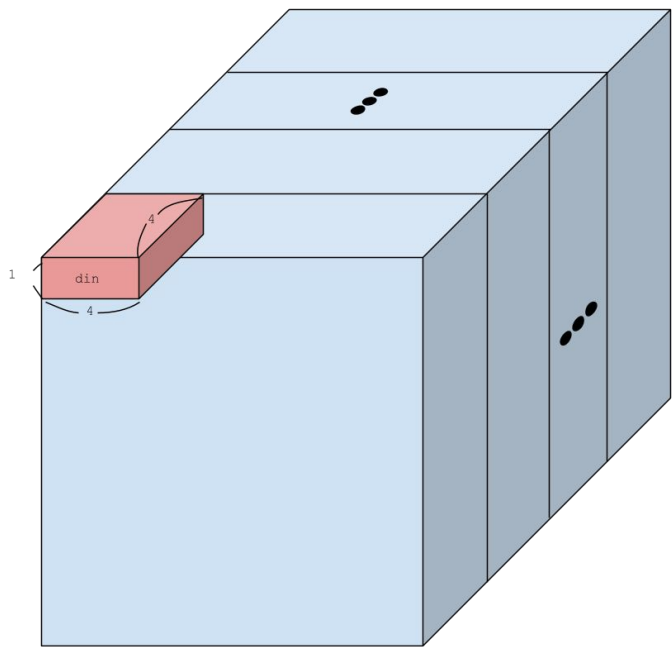
In FIN state, if we get the signal to start reading the reorganized data, move to SEND state



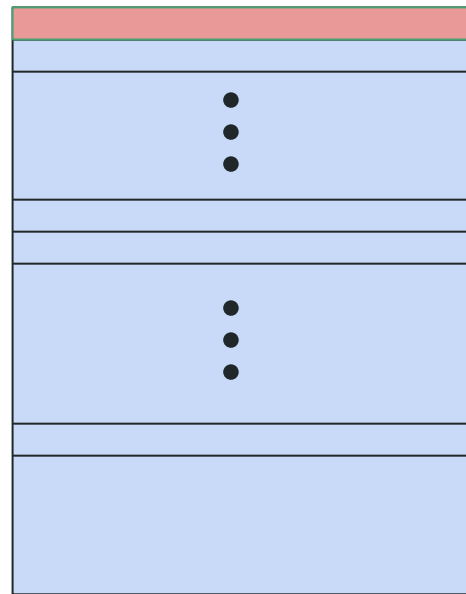
Layer 1 Line Buffer

In SEND state, keep outputting the restructured data until entire row outputted

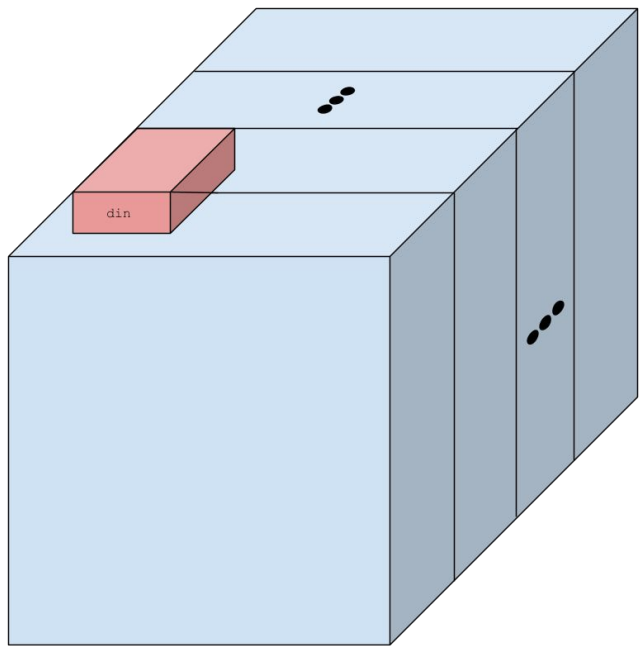




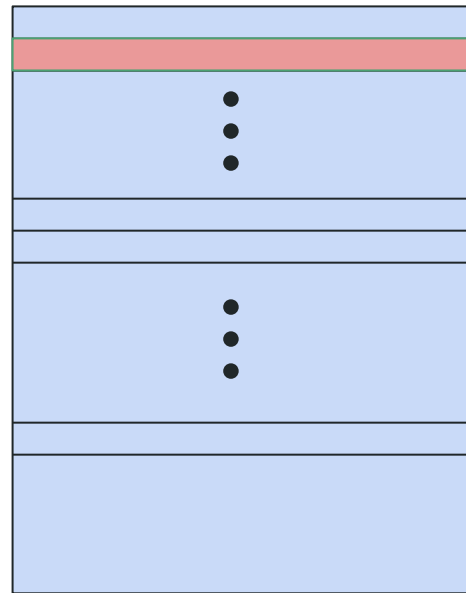
`ch_batch_idx = 0`

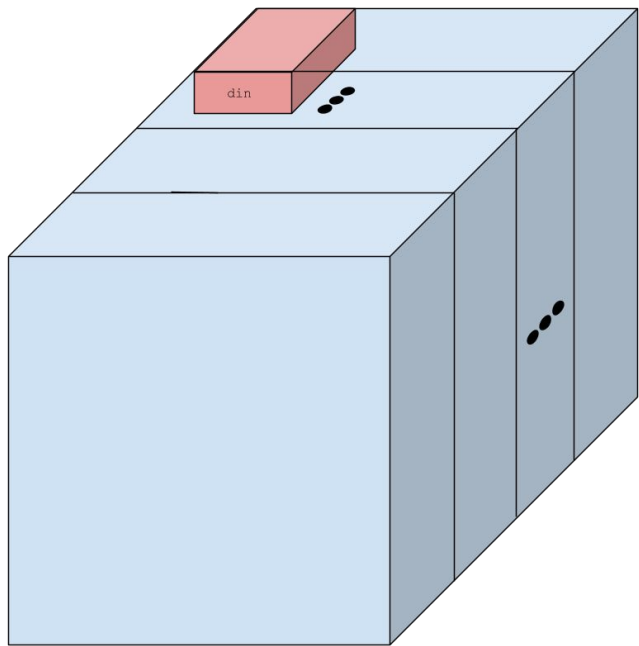


Line buffer FSM in
"output" state

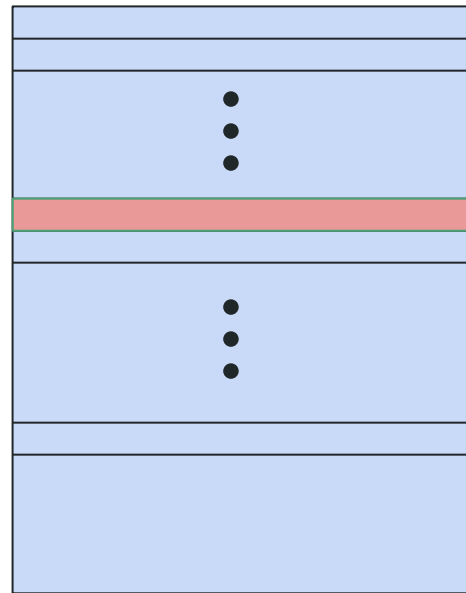


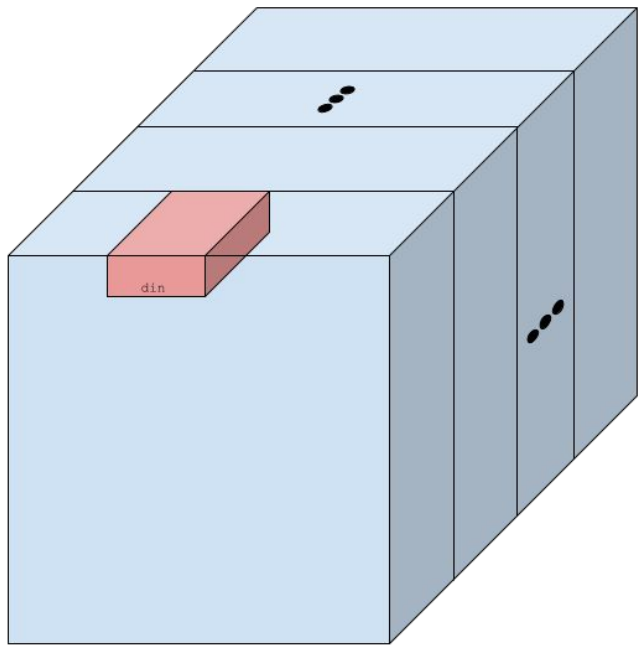
`ch_batch_idx = 1`



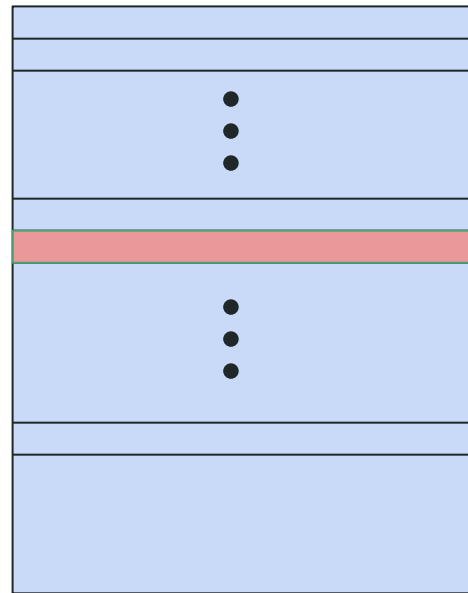


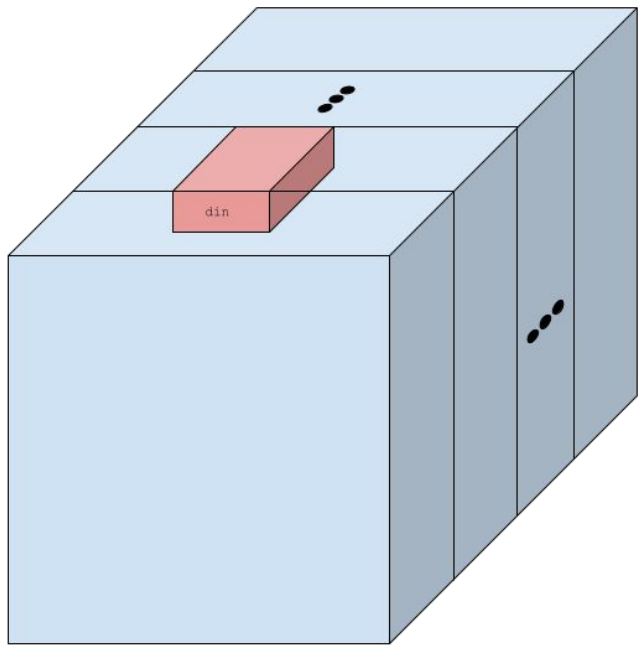
`ch_batch_idx = 7`



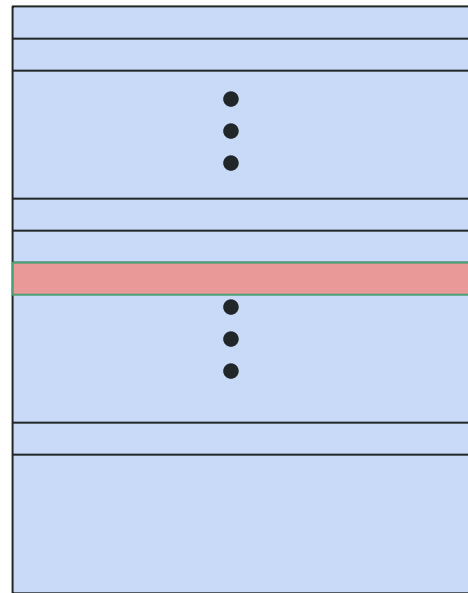


`ch_batch_idx = 0`





`ch_batch_idx = 1`



Layer 2 Convolution Operations

input : $128 \times 128 \times 4$

output : $128 \times 128 \times 4$

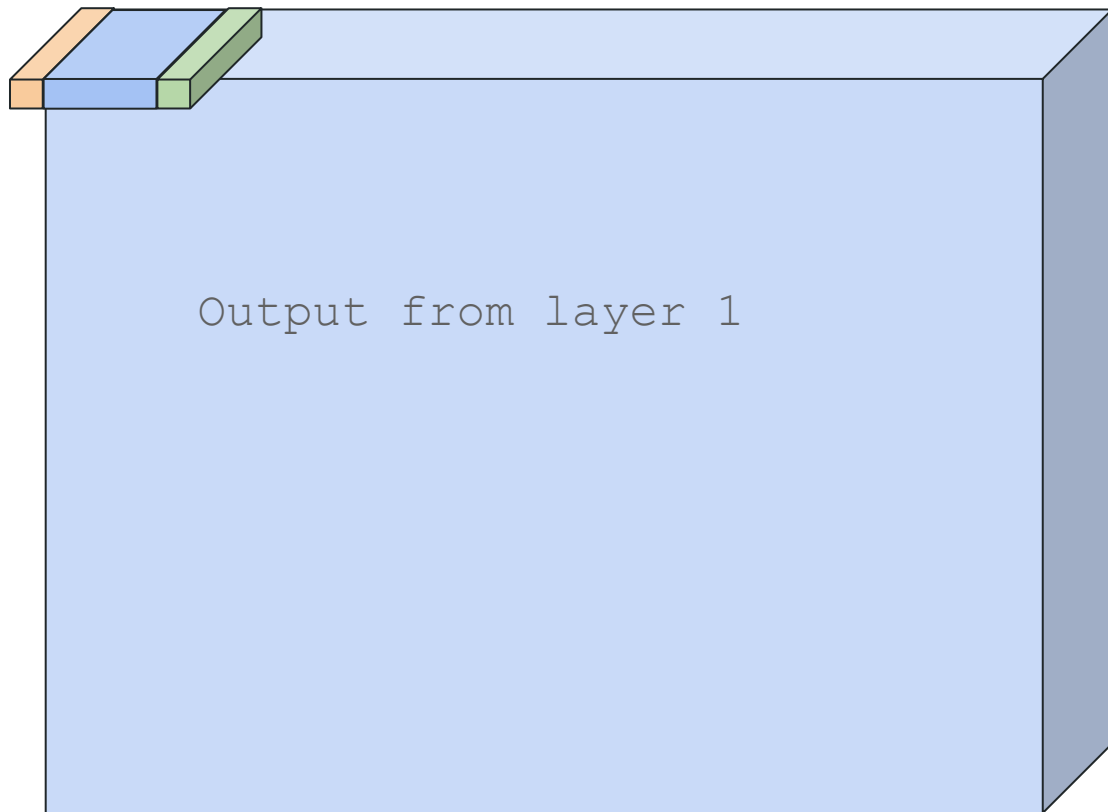
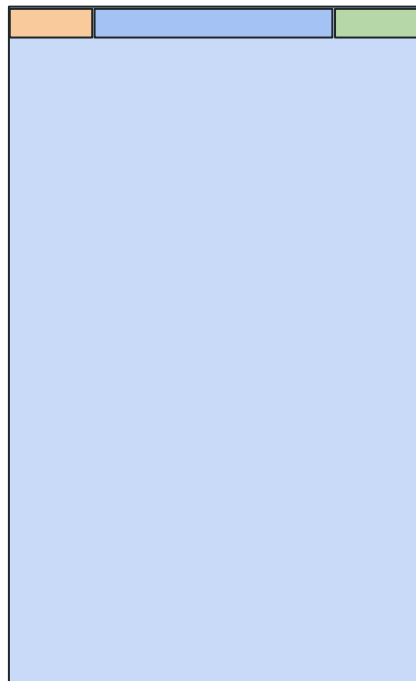
weights : $3 \times 3 \times 4 \times 4$

Now kernel size is 3×3 , so line buffer needs to hold 3 rows (above, current, below)

FSM controls when to increase row, col, etc

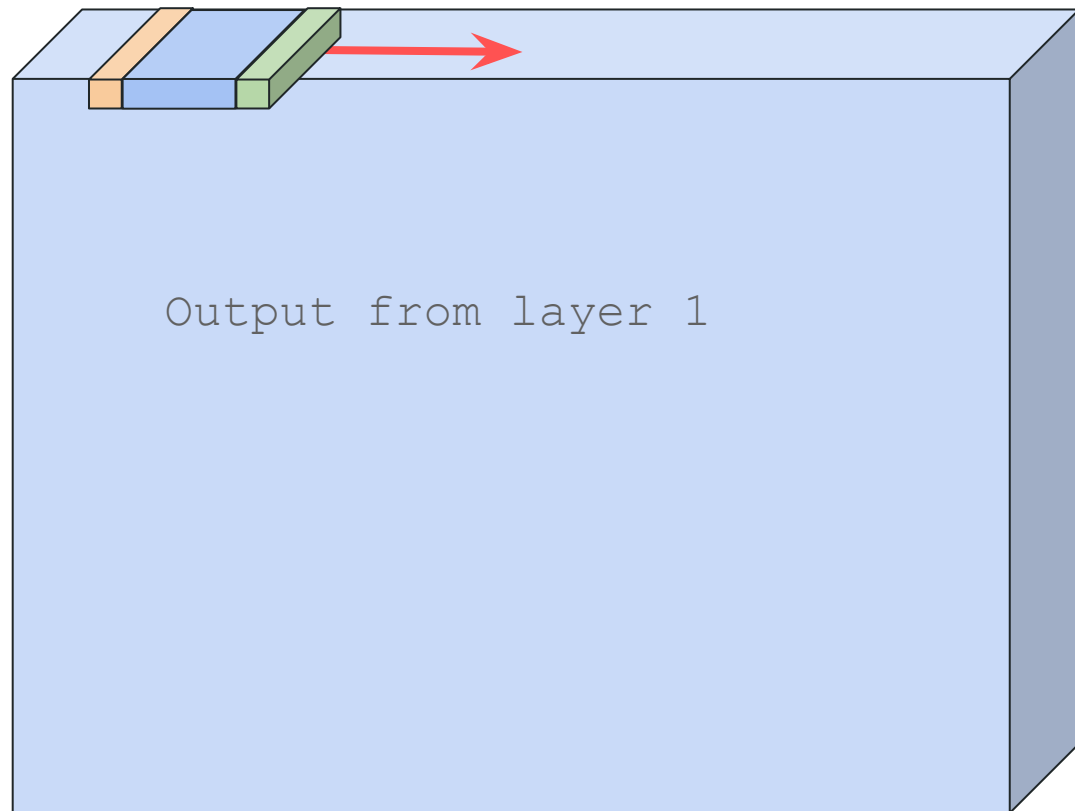
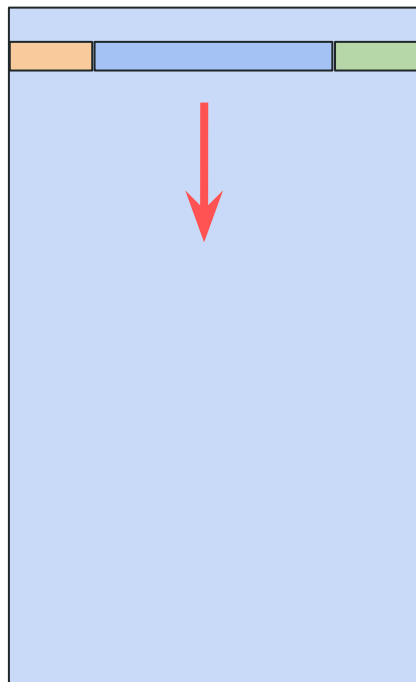
Layer 2 Line Buffer

Layer 2 line buffer



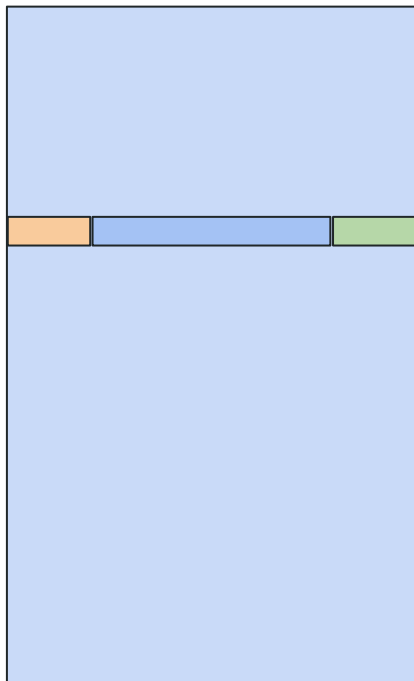
Layer 2 Line Buffer

Layer 2 line buffer

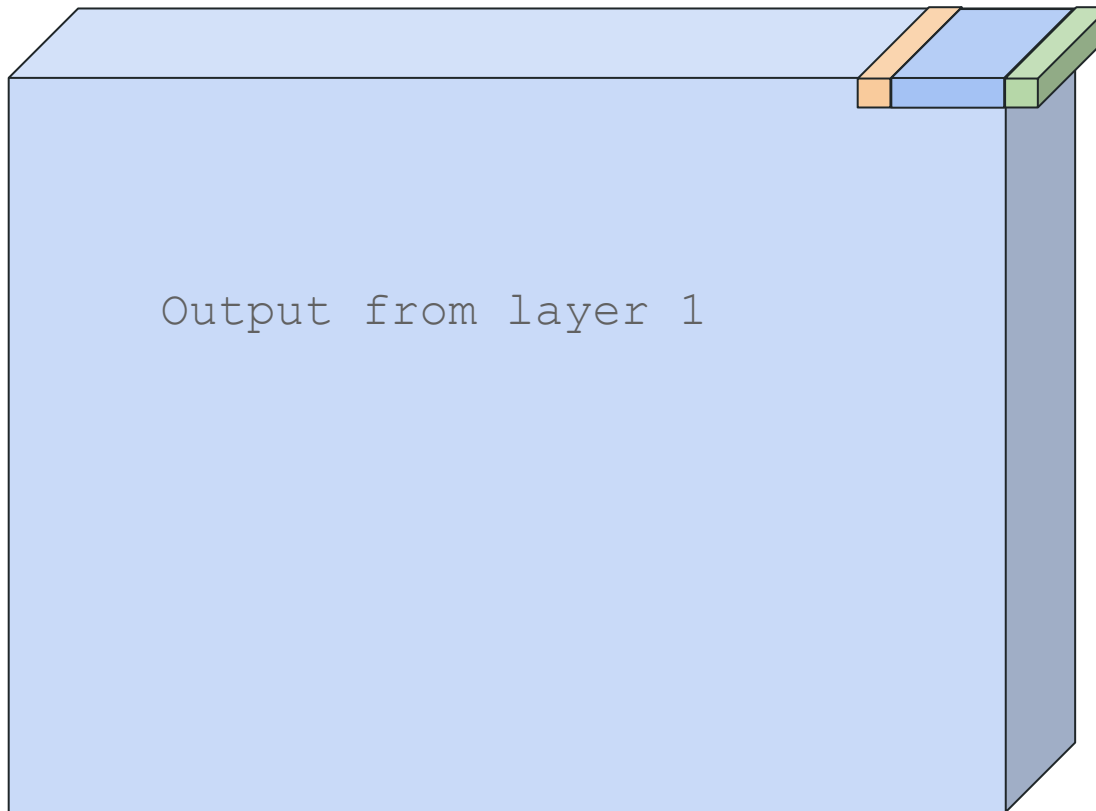


Layer 2 Line Buffer

Layer 2 line buffer

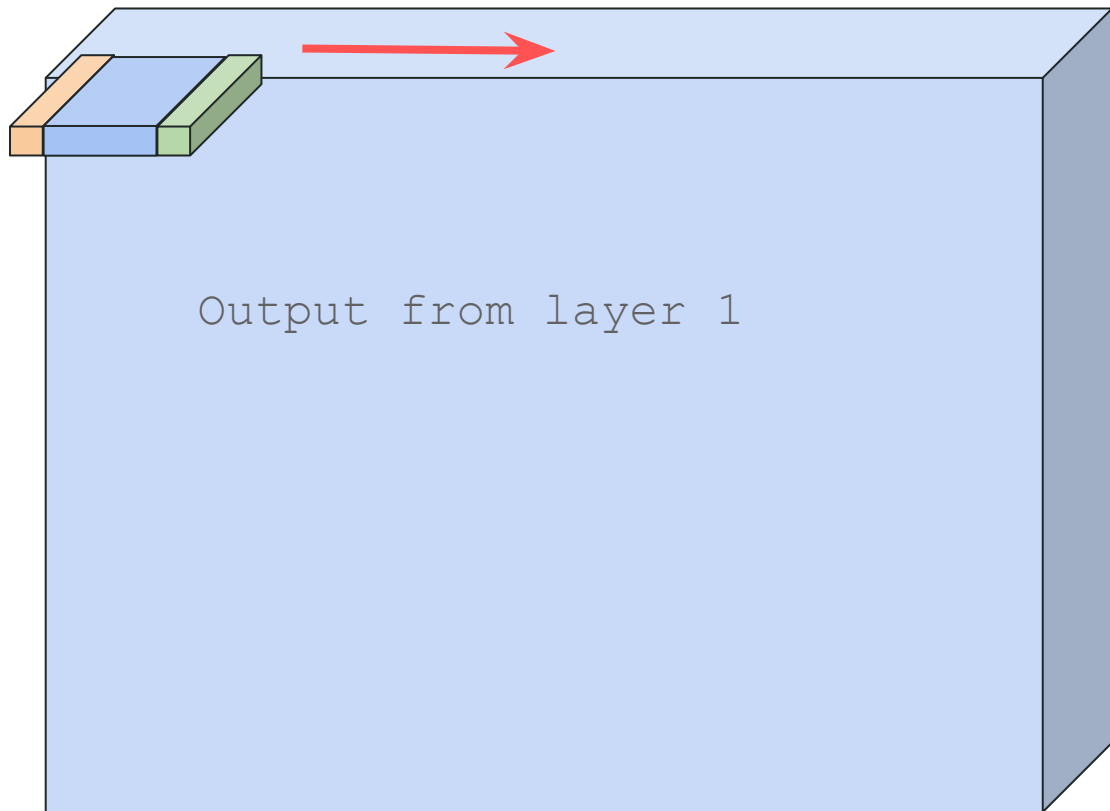
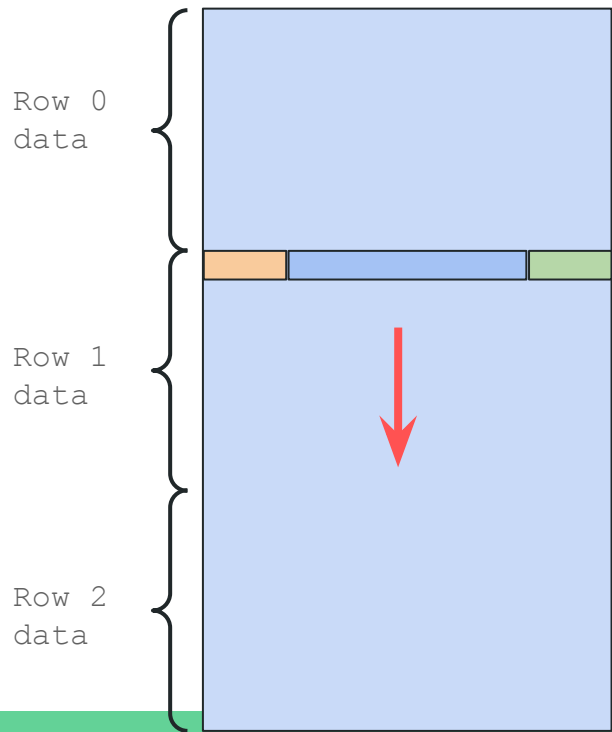


Output from layer 1

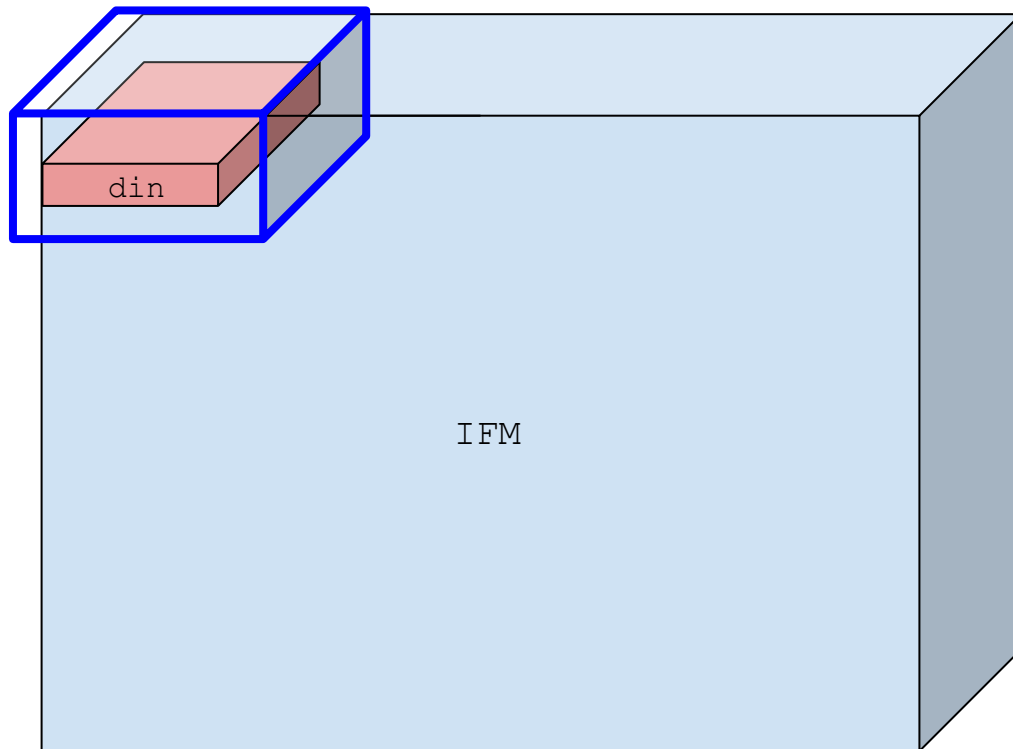


Layer 2 Line Buffer

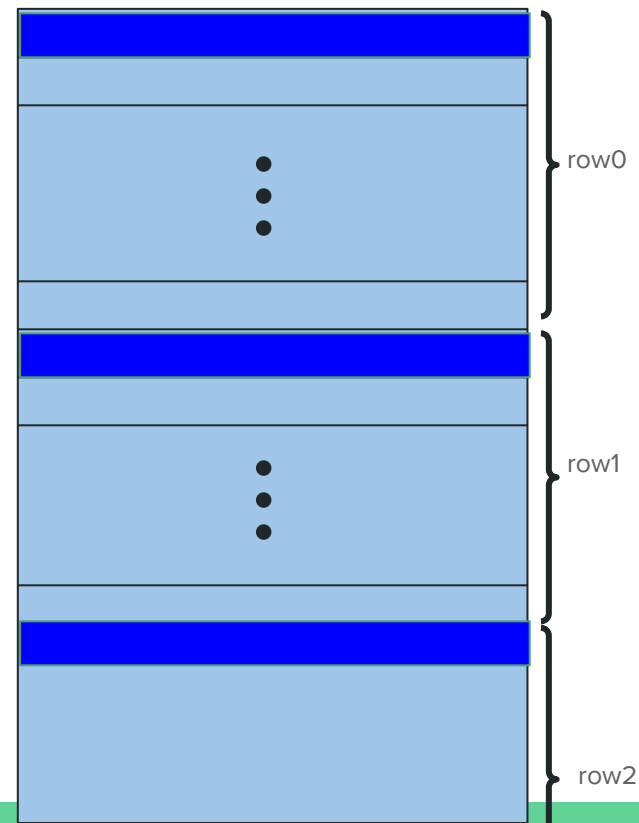
Layer 2 line buffer



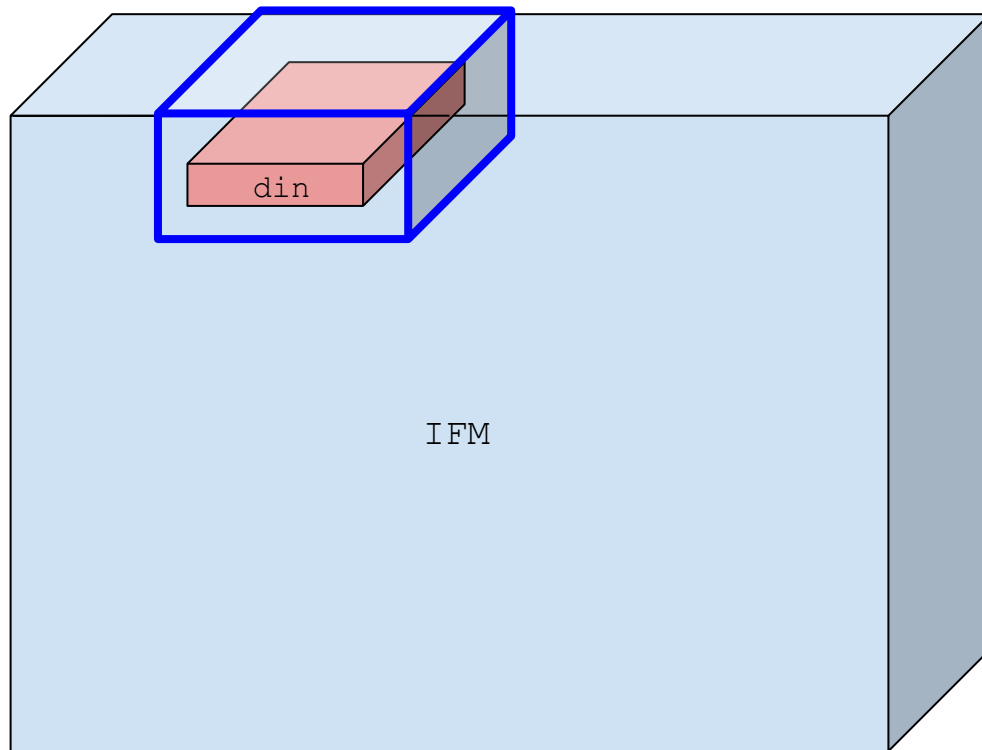
Layer 2 Convolution Operations



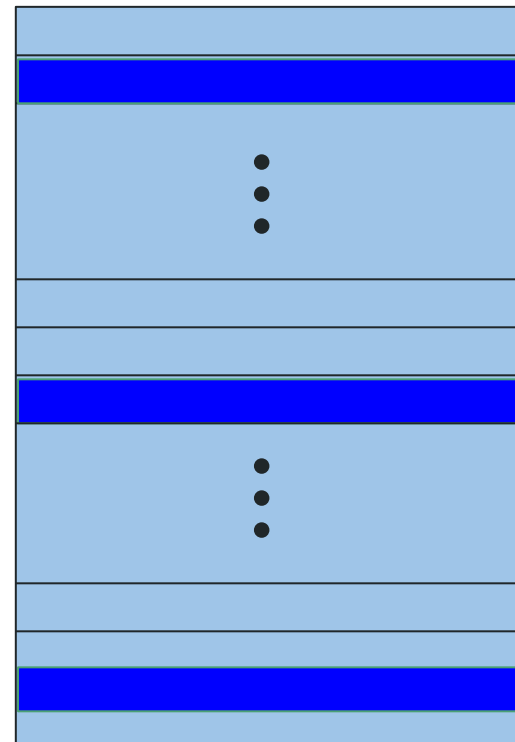
Layer 2 line buffer



Layer 2 Convolution Operations



Layer 2 line buffer



Layer 3 Convolution Operations

input : $128 \times 128 \times 4$

output : $128 \times 128 \times 32$

weights : $1 \times 1 \times 4 \times 32$

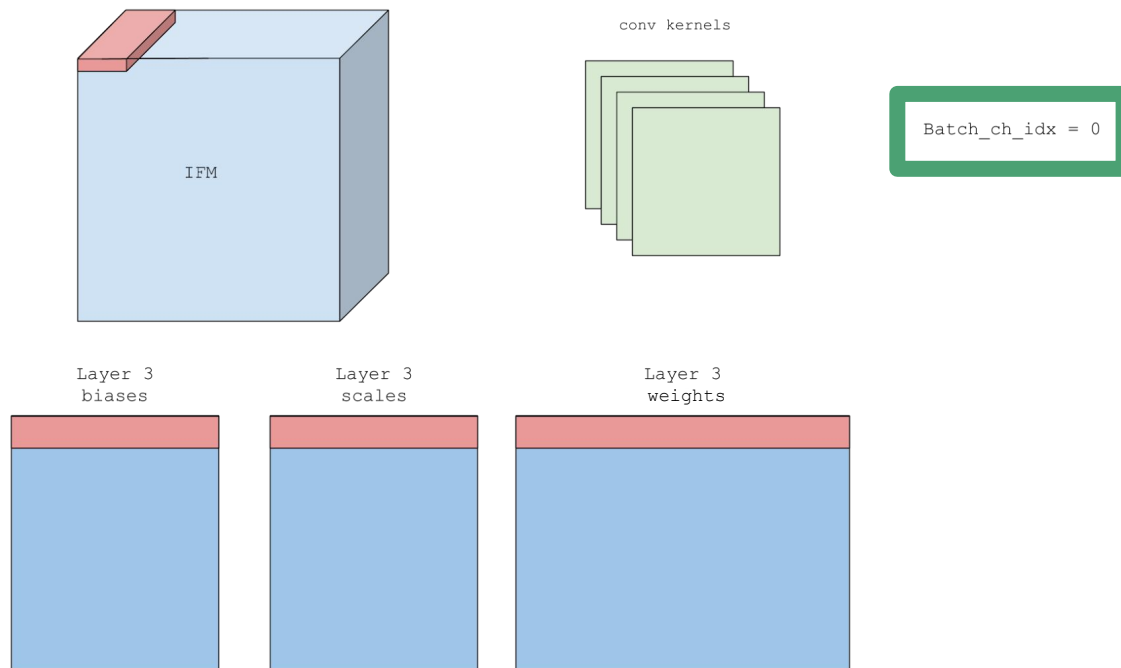
There are 32 filters, so bias and scale must be changed accordingly (in layers 1 and 2, they were static)

batch_ch_idx same as in layer 1, but used to index bias/scales

Fully pipelined accumulate (1 output per clk period)

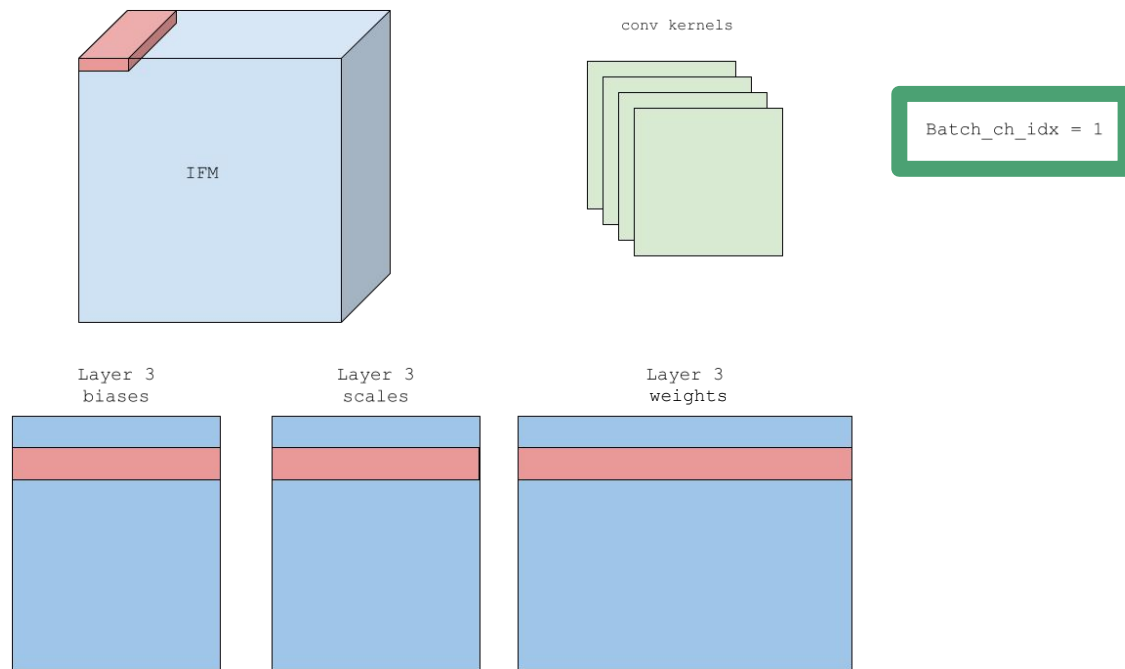
Layer 3 Convolution Operations

Red block = input to kernel



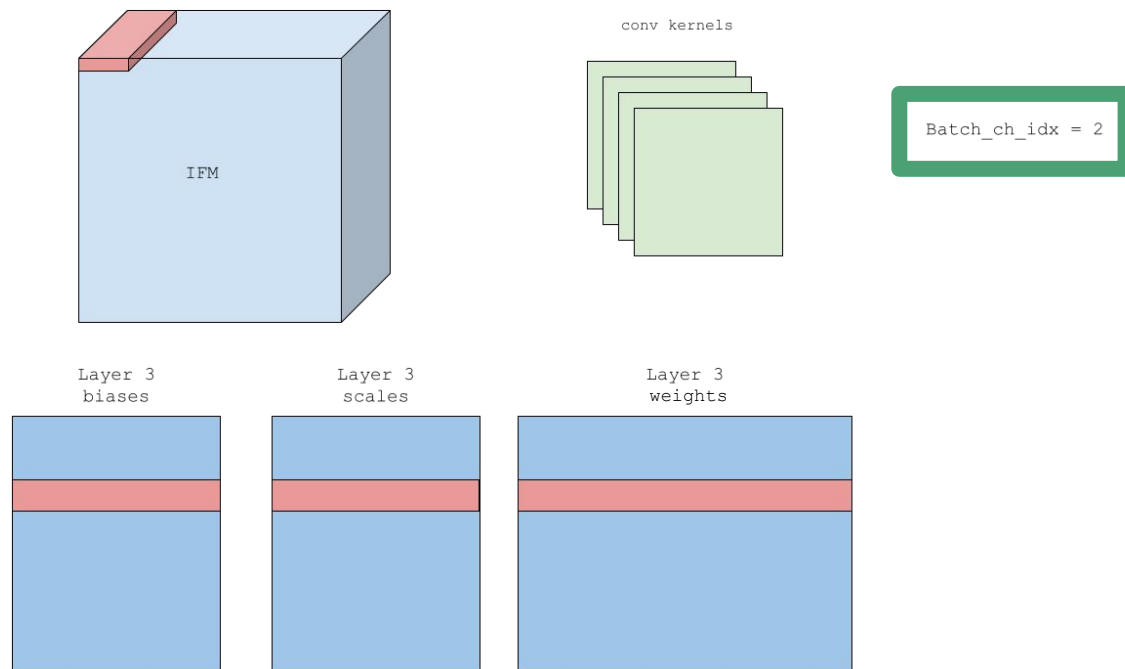
Layer 3 Convolution Operations

Red block = input to kernel



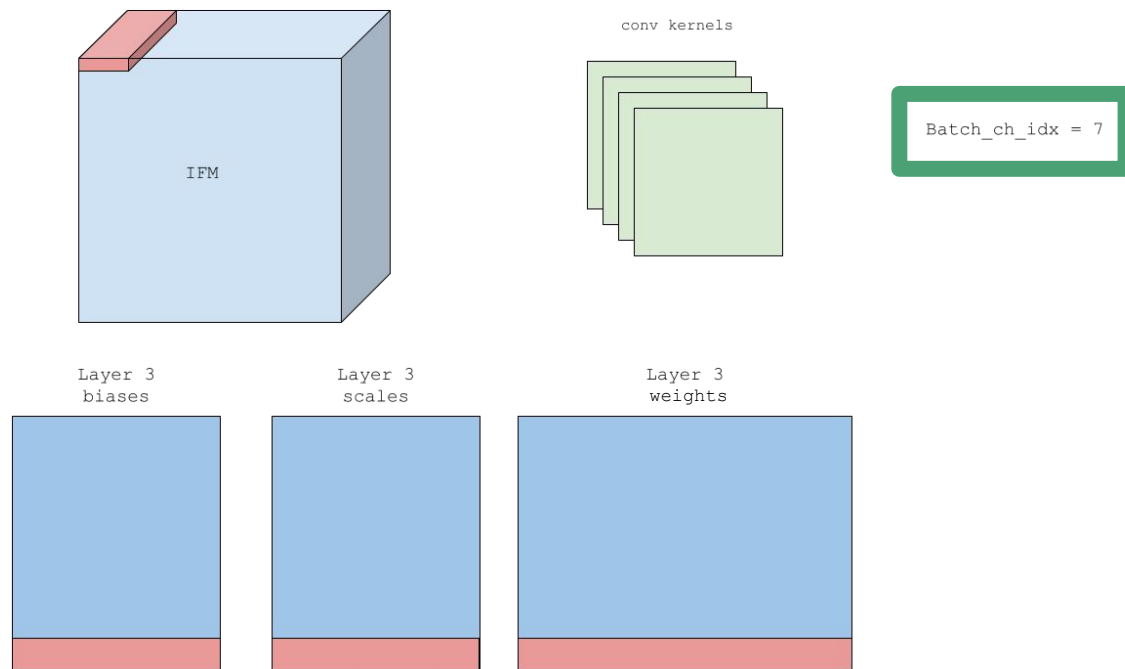
Layer 3 Convolution Operations

Red block = input to kernel

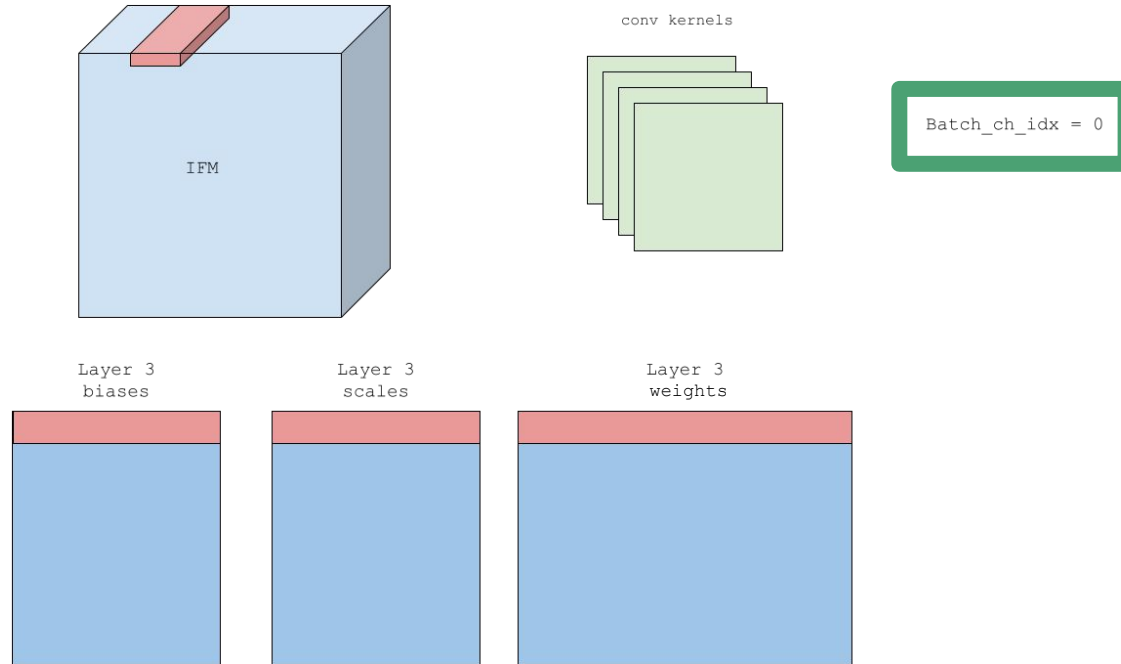


Layer 3 Convolution Operations

Red block = input to kernel



Layer 3 Convolution Operations

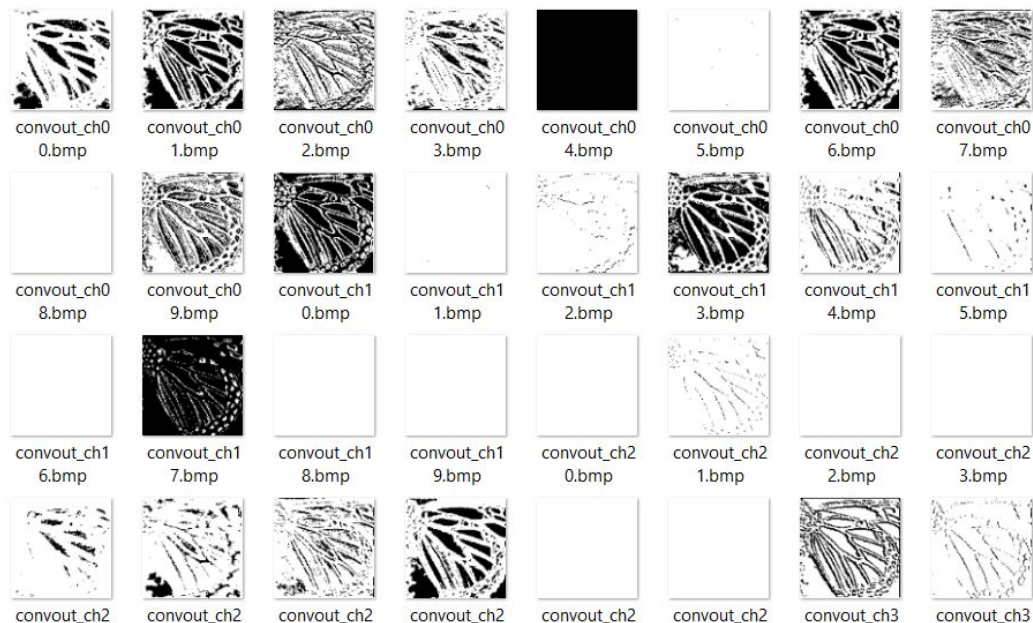


Red block = input to kernel

Increment col/row
and repeat!

Masking

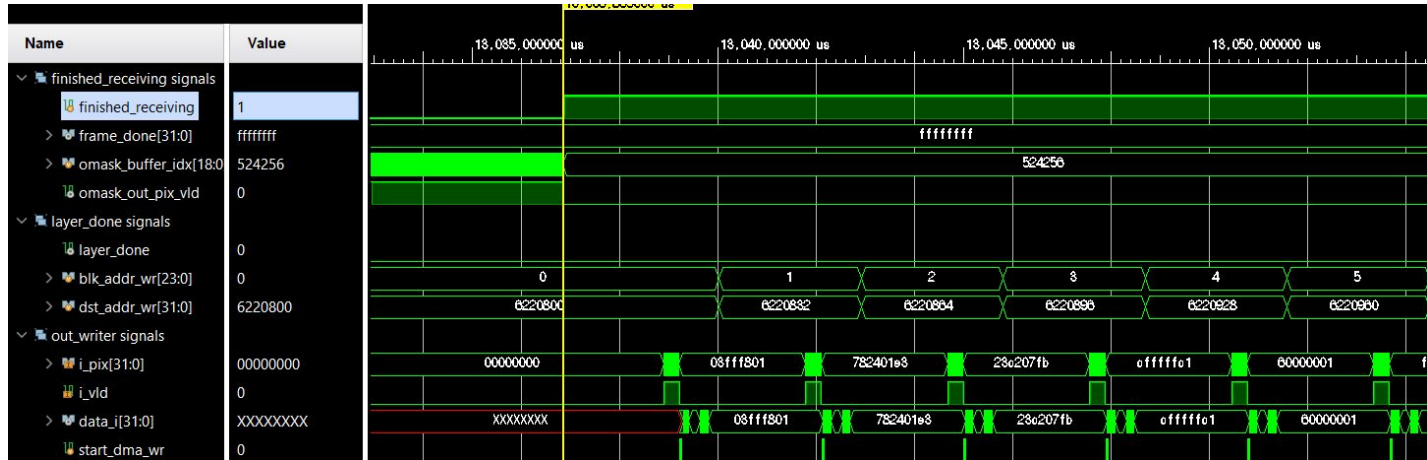
From the third layer, all elements less than or equal to 0 are set to 0, and all elements other than that are set to 255.



Outputting and Writing to Memory

Write mask to memory once all calculations are done

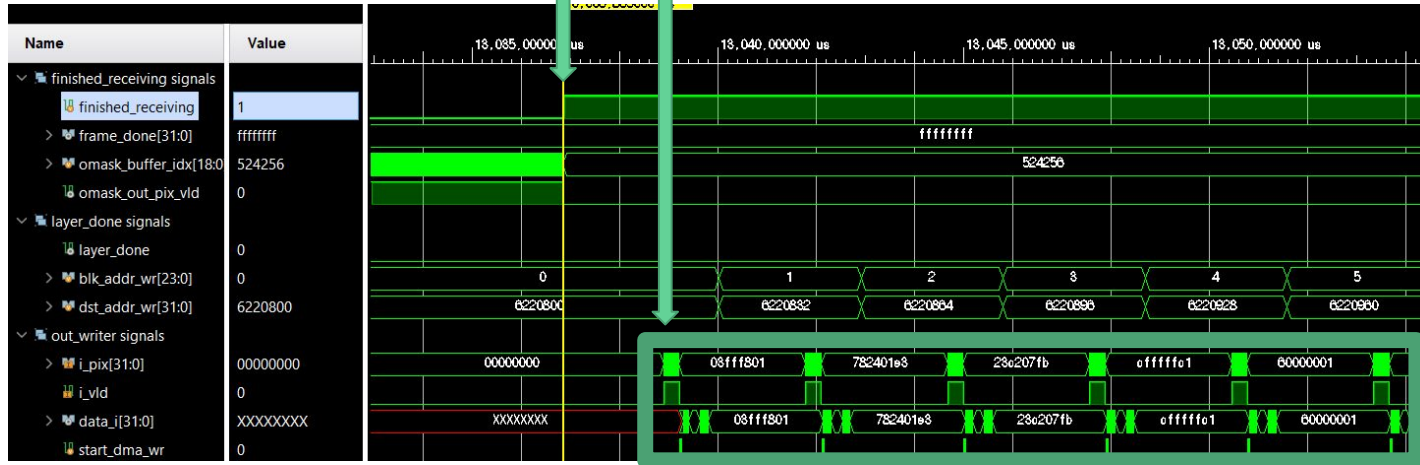
To write to memory, requires bursts of data to be sent at a time



Outputting and Writing to Memory

Finished receiving
signal asserted

Starts sending data in bursts to memory



Checking Results

The outputted OFM's from the last layer are saved as a bmp file and checked against the outputted masks from the matlab SW

```
Results of the channel 02 are same!  
Results of the channel 03 are same!  
Results of the channel 04 are same!  
Results of the channel 05 are same!  
Results of the channel 06 are same!  
Results of the channel 07 are same!  
Results of the channel 08 are same!  
Results of the channel 09 are same!  
Results of the channel 10 are same!  
Results of the channel 11 are same!  
Results of the channel 12 are same!  
Results of the channel 13 are same!  
Results of the channel 14 are same!  
Results of the channel 15 are same!  
Results of the channel 16 are same!  
Results of the channel 17 are same!  
Results of the channel 18 are same!  
Results of the channel 19 are same!  
Results of the channel 20 are same!  
Results of the channel 21 are same!  
Results of the channel 22 are same!  
Results of the channel 23 are same!  
Results of the channel 24 are same!  
Results of the channel 25 are same!  
Results of the channel 26 are same!  
Results of the channel 27 are same!  
Results of the channel 28 are same!  
Results of the channel 29 are same!  
Results of the channel 30 are same!  
Results of the channel 31 are same!  
Results of the channel 32 are same!
```

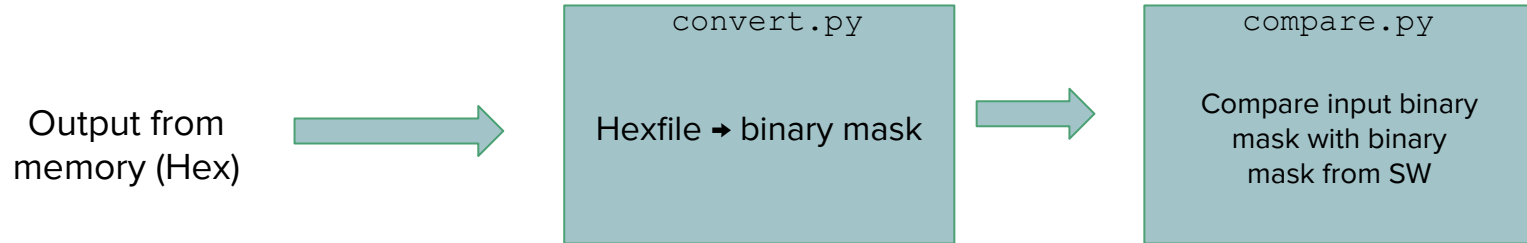
Checking Results

Outputs are verified for all layers (3,4,5, and 6)

```
Results of the channel 02 are same!  
Results of the channel 03 are same!  
Results of the channel 04 are same!  
Results of the channel 05 are same!  
Results of the channel 06 are same!  
Results of the channel 07 are same!  
Results of the channel 08 are same!  
Results of the channel 09 are same!  
Results of the channel 10 are same!  
Results of the channel 11 are same!  
Results of the channel 12 are same!  
Results of the channel 13 are same!  
Results of the channel 14 are same!  
Results of the channel 15 are same!  
Results of the channel 16 are same!  
Results of the channel 17 are same!  
Results of the channel 18 are same!  
Results of the channel 19 are same!  
Results of the channel 20 are same!  
Results of the channel 21 are same!  
Results of the channel 22 are same!  
Results of the channel 23 are same!  
Results of the channel 24 are same!  
Results of the channel 25 are same!  
Results of the channel 26 are same!  
Results of the channel 27 are same!  
Results of the channel 28 are same!  
Results of the channel 29 are same!  
Results of the channel 30 are same!  
Results of the channel 31 are same!  
Results of the channel 32 are same!
```

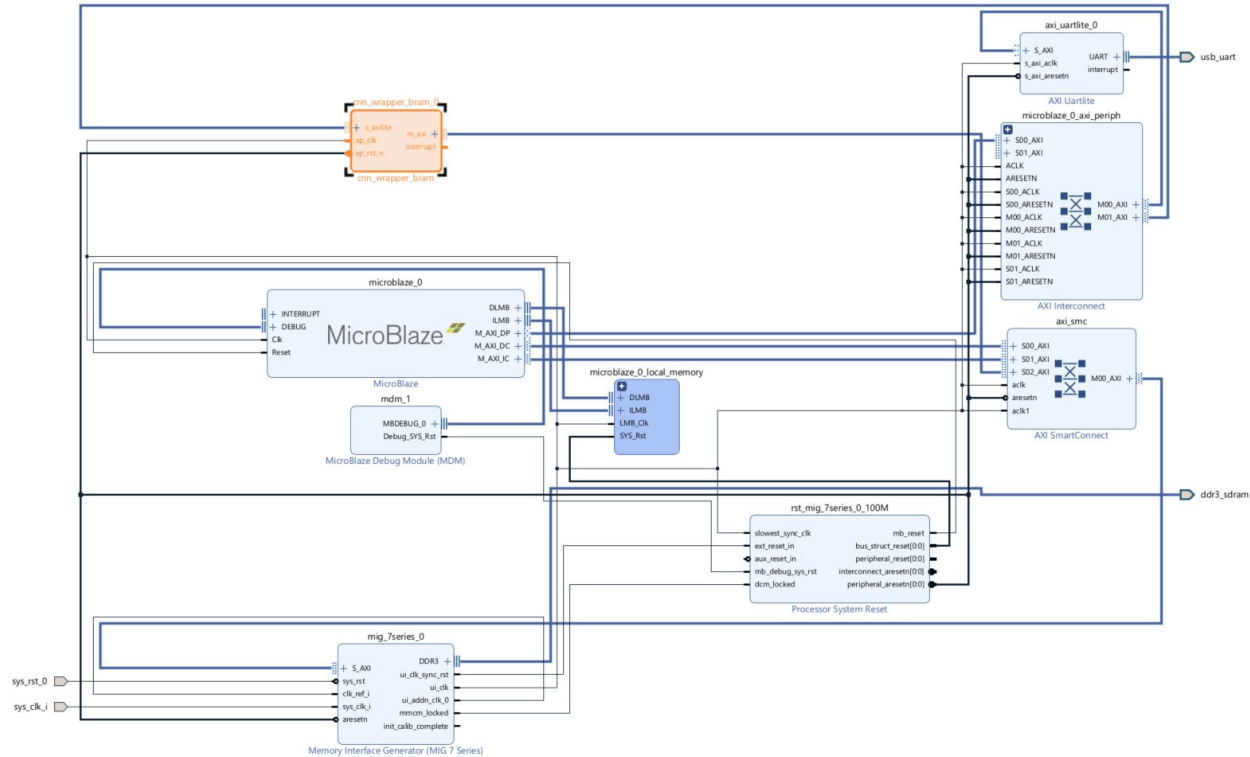

Checking results (extra, just in case)

Output converted to binary, and compared with SW output



The two files are identical.

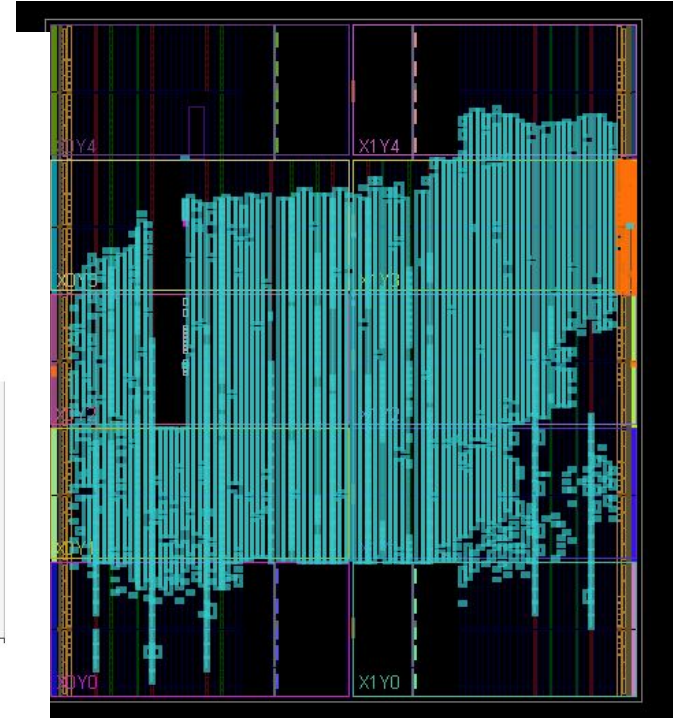
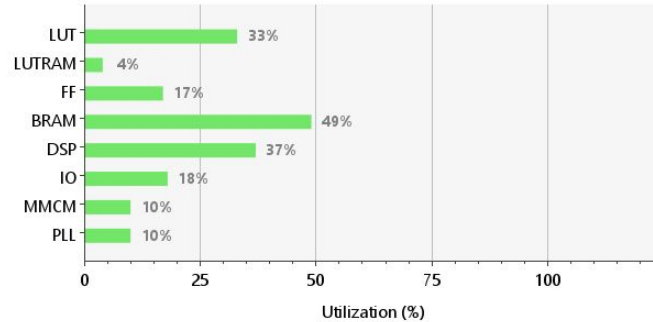
Synthesized and Implemented Design



Synthesized and Implemented Design

Nexys Video
FPGA board

Resource	Utilization	Available	Utilization %
LUT	44535	134600	33.09
LUTRAM	1827	46200	3.95
FF	44879	269200	16.67
BRAM	180	365	49.32
DSP	272	740	36.76
IO	52	285	18.25
MMCM	1	10	10.00
PLL	1	10	10.00



TODO

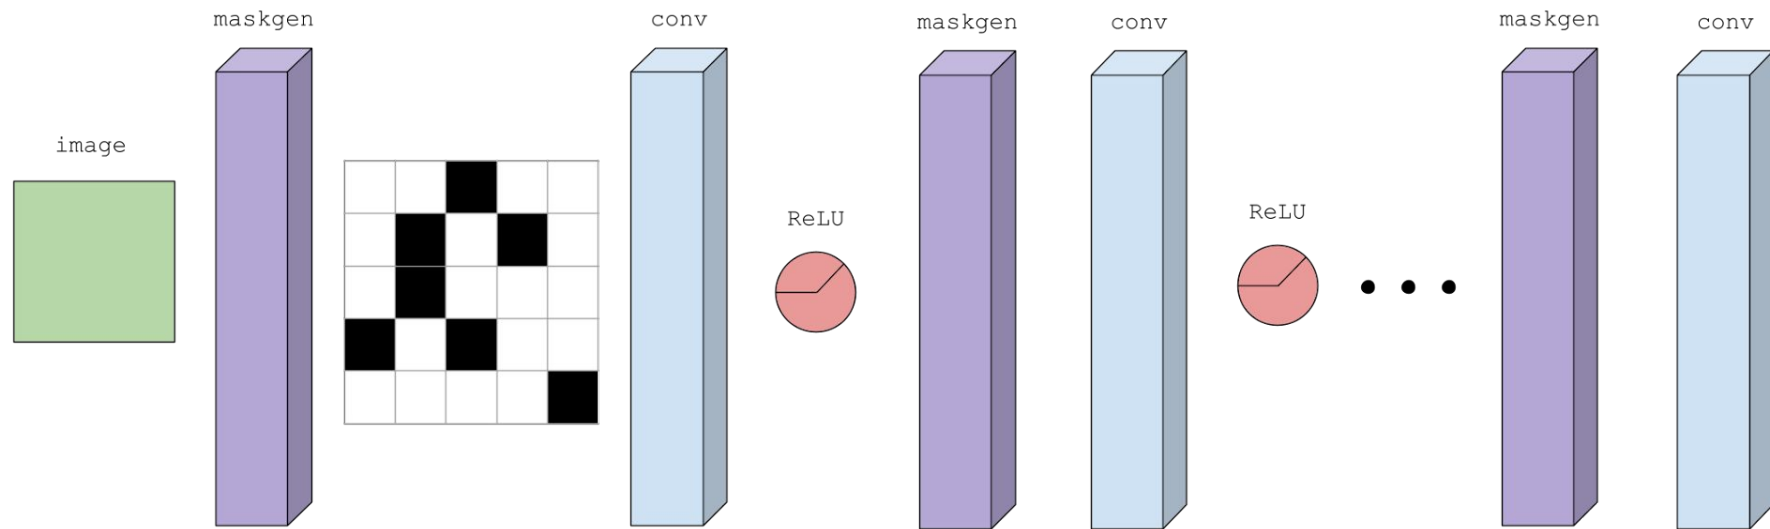
Get it running on an FPGA...

```
[ OK ] MODE 0x01
[ OK ] Response: Hello, World!
sending width = 0x080, height = 0x080
[ OK ] MODE 0x05
[ OK ] Response: Waiting_CMD
sending data to : 0x14
data to send is : 0x3080080
[ OK ] SEND CMD
[ OK ] Response Store complete
checking width, height
[ OK ] MODE 0x06
[ OK ] Response: Waiting_CMD
[ OK ] SEND OFFSET
[ OK ] Response: 0x00000080
```

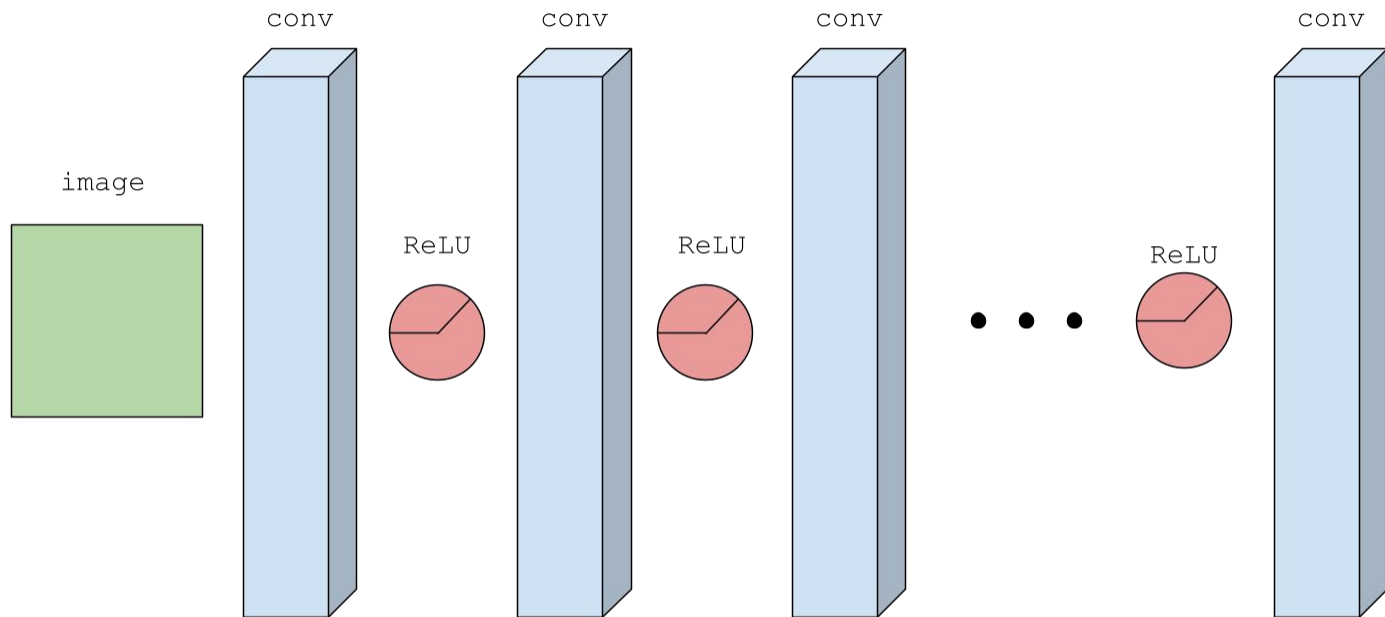


Extra slides

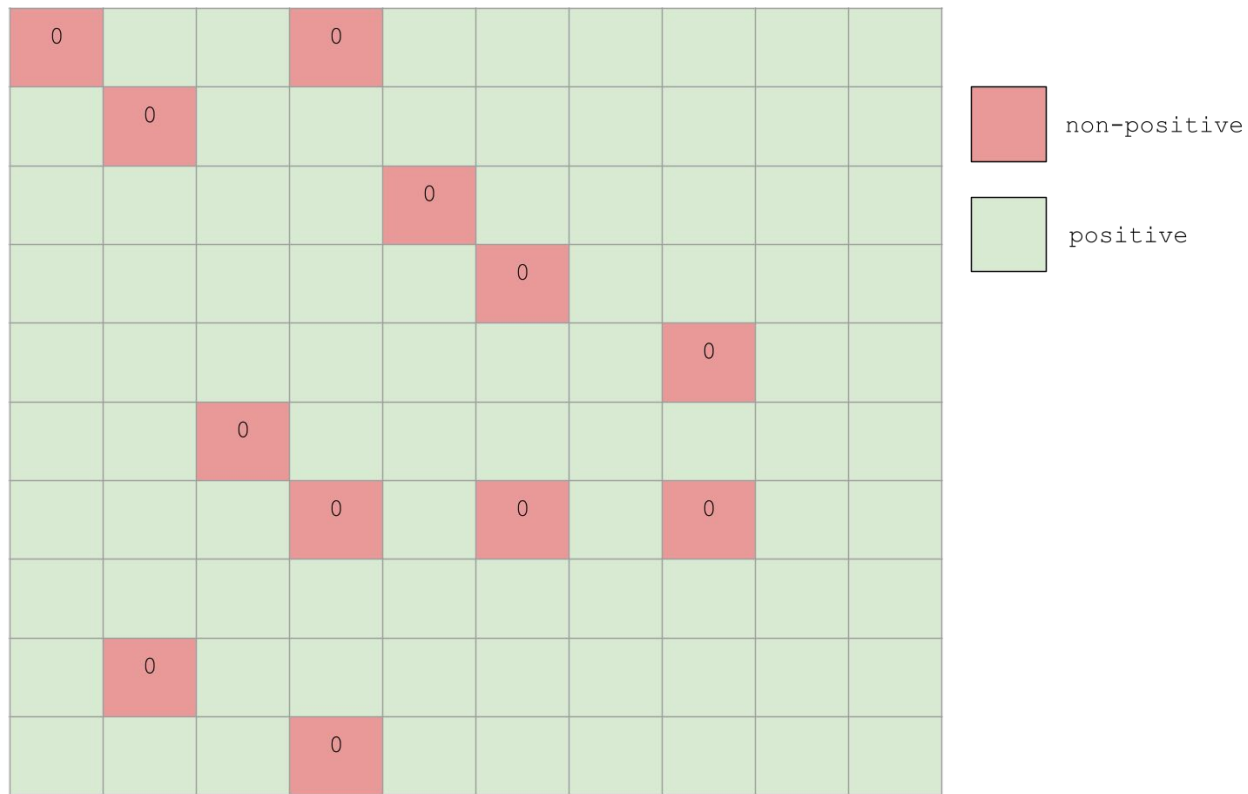
Zero Skipping via Mask Generation - The Big Picture



Why Zero Skipping?

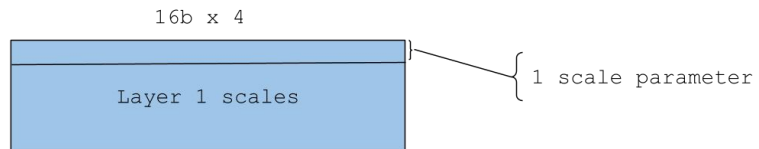
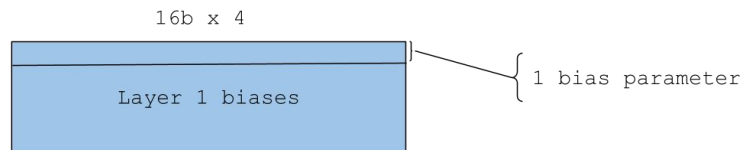
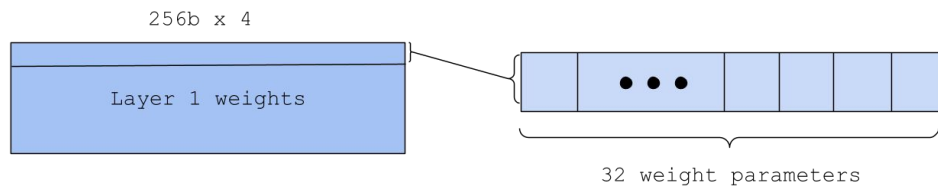


After ReLU



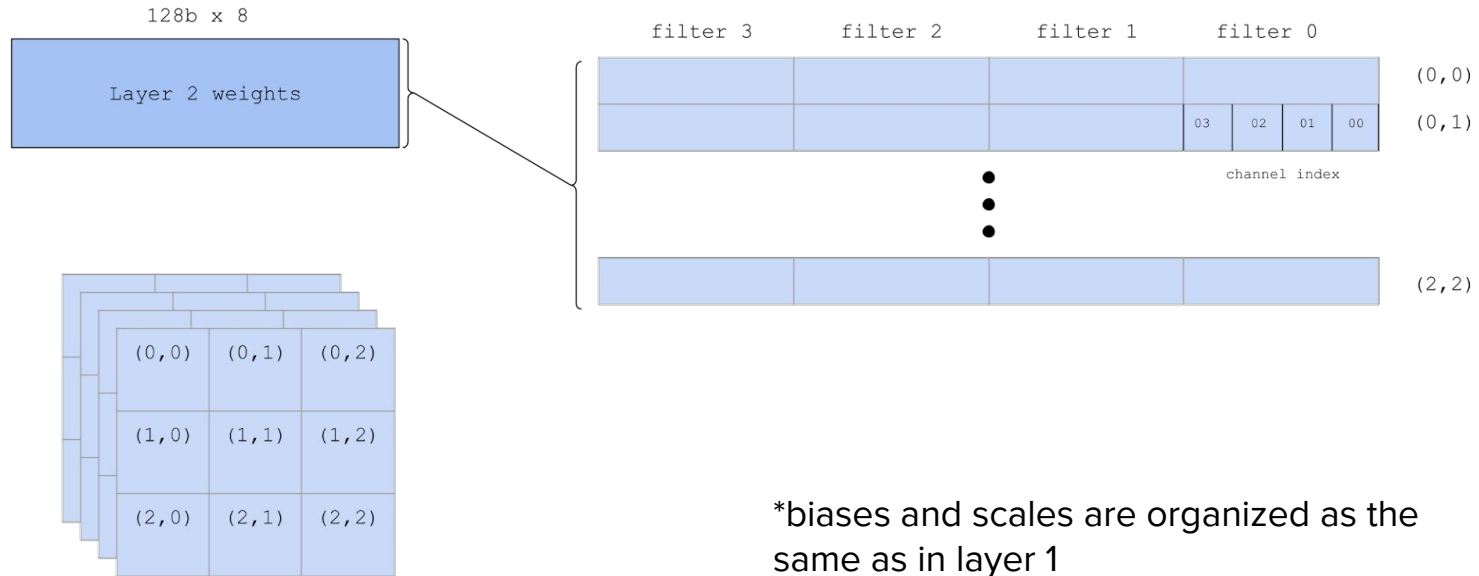
Inputs

Layer 1 conv weights are 1x1x32x4



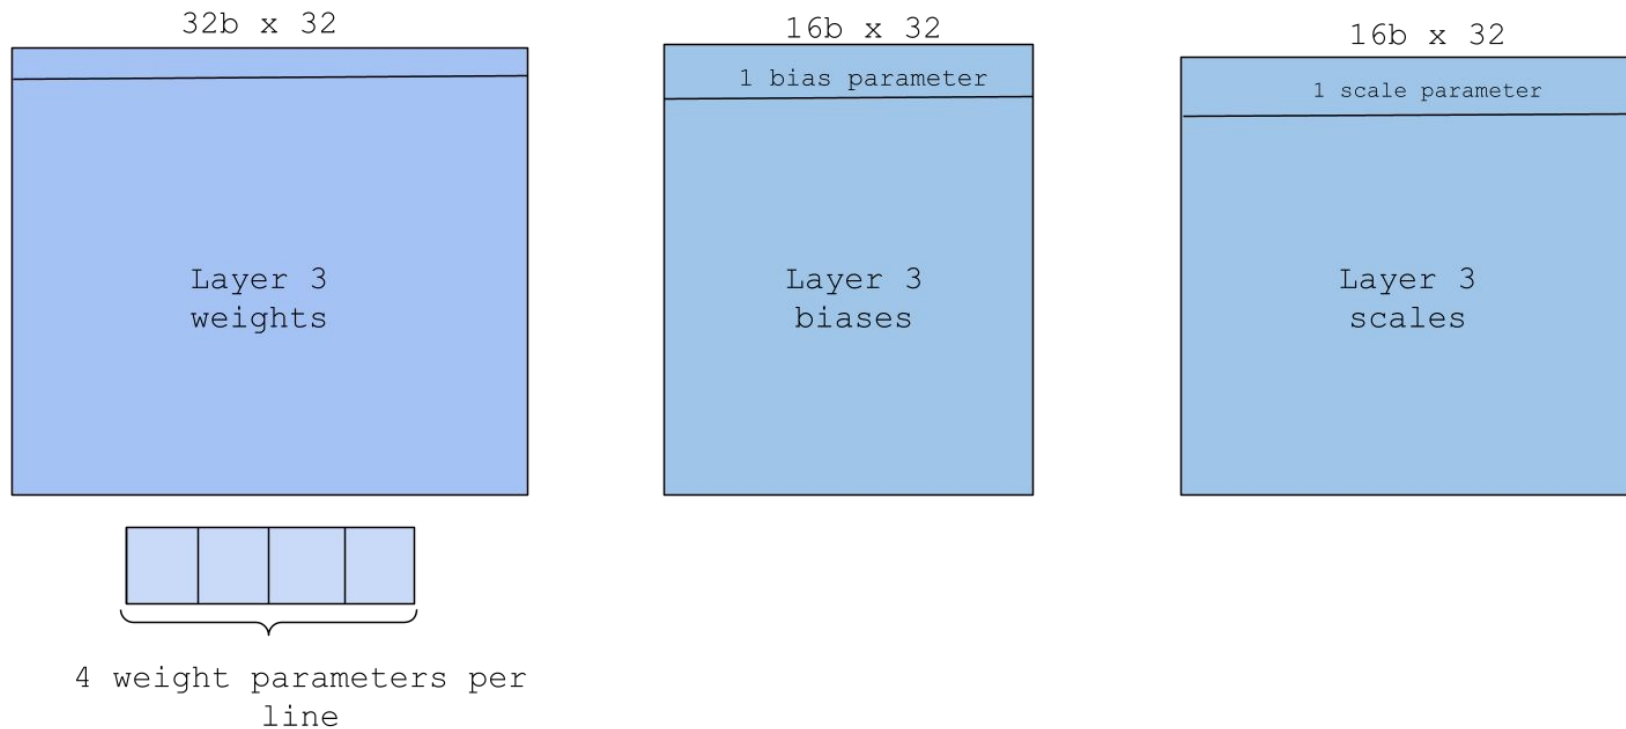
Inputs

Layer 2 conv weights are 3x3x4x4












Inputs

Layer 3 conv weights are 1x1x4x32



Initializing Memory with ROMs

```
bias_blk_mem u_bias_blk_mem(  
    //input  */clk      (clk      ),  
    //input  */addr     (bias_blk_mem_addr ),  
    //input  */ena      (bias_blk_mem_ena  ),  
    //output */douta    (bias_blk_mem_data ),  
);  
  
reg [7:0] scale_blk_mem_addr;  
//reg [7:0] scale_blk_mem_addr_d;  
reg      scale_blk_mem_ena;  
wire [15:0] scale_blk_mem_data;  
scale_blk_mem u_scale_blk_mem(  
    //input  */clk      (clk      ),  
    //input  */addr     (scale_blk_mem_addr ),  
    //input  */ena      (scale_blk_mem_ena  ),  
    //output */douta    (scale_blk_mem_data ),  
);  
  
reg [4:0] weight_0_blk_mem_addr;  
//reg [4:0] weight_0_blk_mem_addr_d;  
reg      weight_0_blk_mem_ena;  
wire [255:0] weight_0_blk_mem_data;  
weight_0_blk_mem u_weight_0_blk_mem(  
    //input  */clk      (clk      ),  
    //input  */addr     (weight_0_blk_mem_addr ),  
    //input  */ena      (weight_0_blk_mem_ena  ),  
    //output */douta    (weight_0_blk_mem_data ),  
);  
  
reg [5:0] weight_1_blk_mem_addr;  
//reg [5:0] weight_1_blk_mem_addr_d;  
reg      weight_1_blk_mem_ena;  
wire [127:0] weight_1_blk_mem_data;  
weight_1_blk_mem u_weight_1_blk_mem(  
    //input  */clk      (clk      ),  
    //input  */addr     (weight_1_blk_mem_addr ),  
    //input  */ena      (weight_1_blk_mem_ena  ),  
    //output */douta    (weight_1_blk_mem_data ),  
);  
  
reg [7:0] weight_2_blk_mem_addr;  
//reg [7:0] weight_2_blk_mem_addr_d;  
reg      weight_2_blk_mem_ena;  
wire [31:0] weight_2_blk_mem_data;  
weight_2_blk_mem u_weight_2_blk_mem(  
    //input  */clk      (clk      ),  
    //input  */addr     (weight_2_blk_mem_addr ),  
    //input  */ena      (weight_2_blk_mem_ena  ),  
    //output */douta    (weight_2_blk_mem_data ),  
);
```

- >   u_bias_blk_mem : bias_blk_mem (bias_blk_mem.xci)
- >   u_scale_blk_mem : scale_blk_mem (scale_blk_mem.xci)
- >   u_weight_0_blk_mem : weight_0_blk_mem (weight_0_blk_mem.xci)
- >   u_weight_1_blk_mem : weight_1_blk_mem (weight_1_blk_mem.xci)
- >   u_weight_2_blk_mem : weight_2_blk_mem (weight_2_blk_mem.xci)

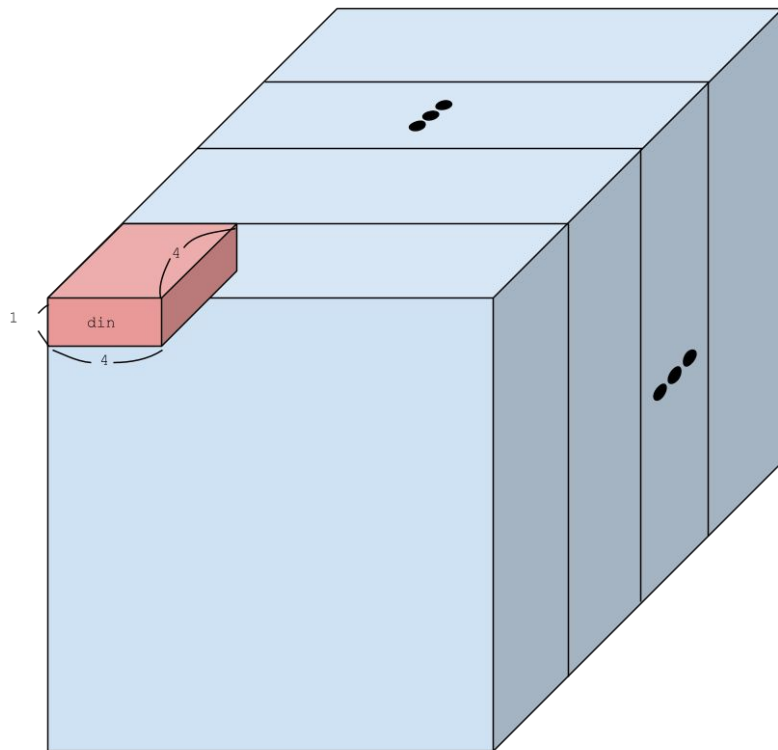
5 ROMs

1 for all scales, biases, weights

3 for all weights for each layer

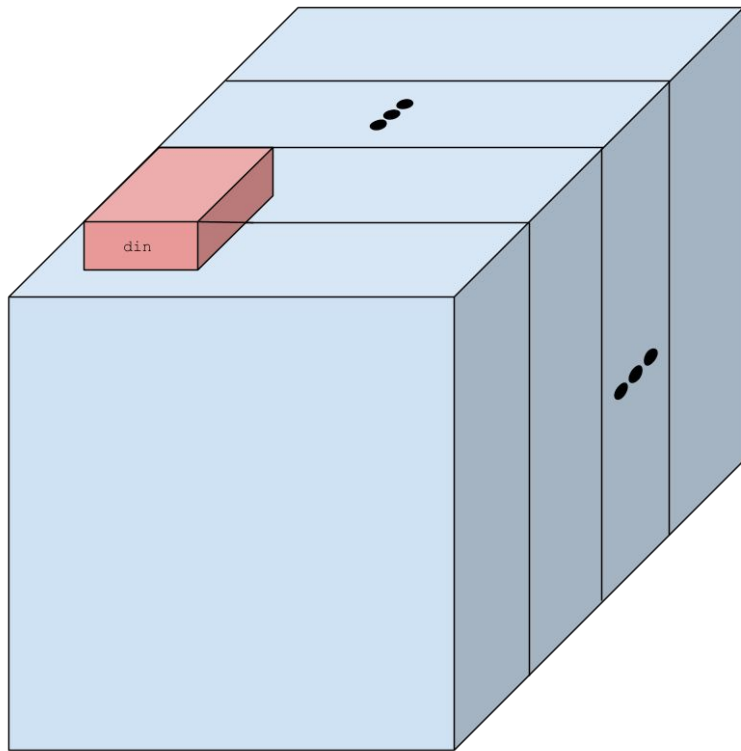
See extra slides for detailed file structure

Layer 1 Convolution Operations



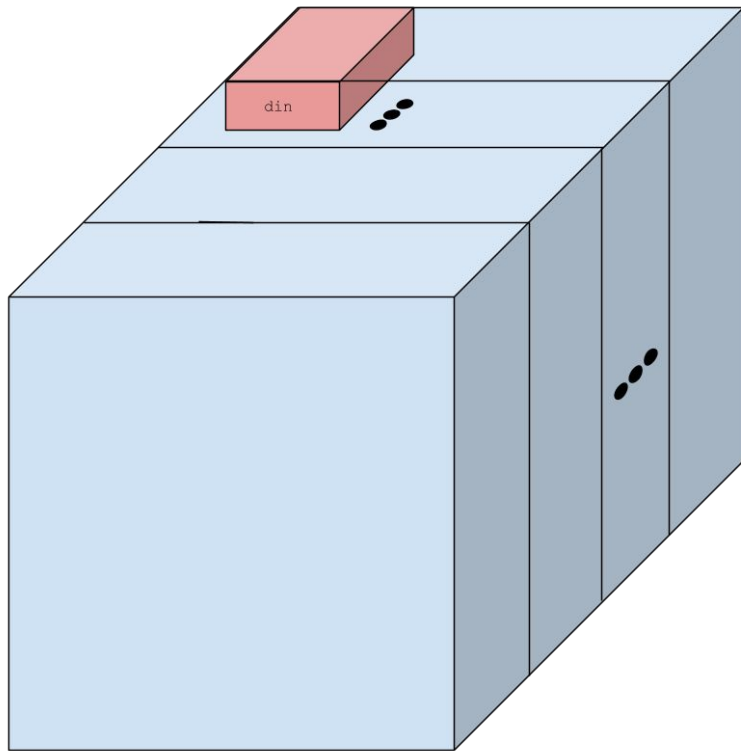
`ch_batch_idx = 0`

Layer 1 Convolution Operations



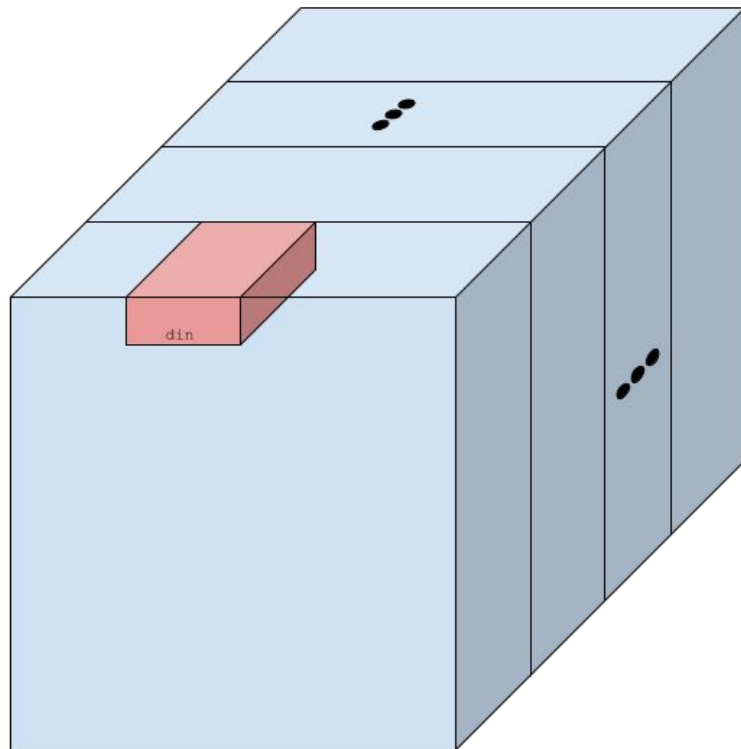
`ch_batch_idx = 1`

Layer 1 Convolution Operations



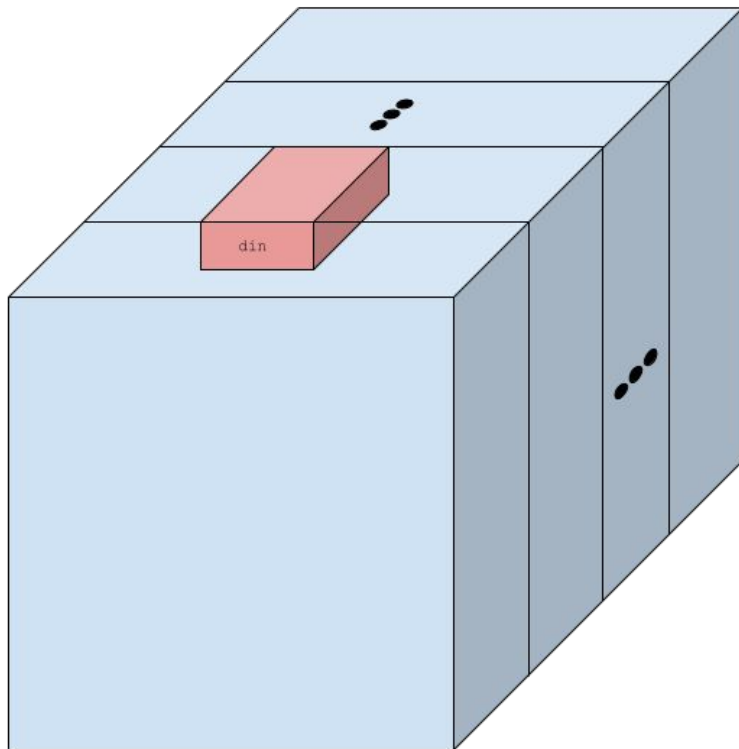
`ch_batch_idx = 7`

Layer 1 Convolution Operations



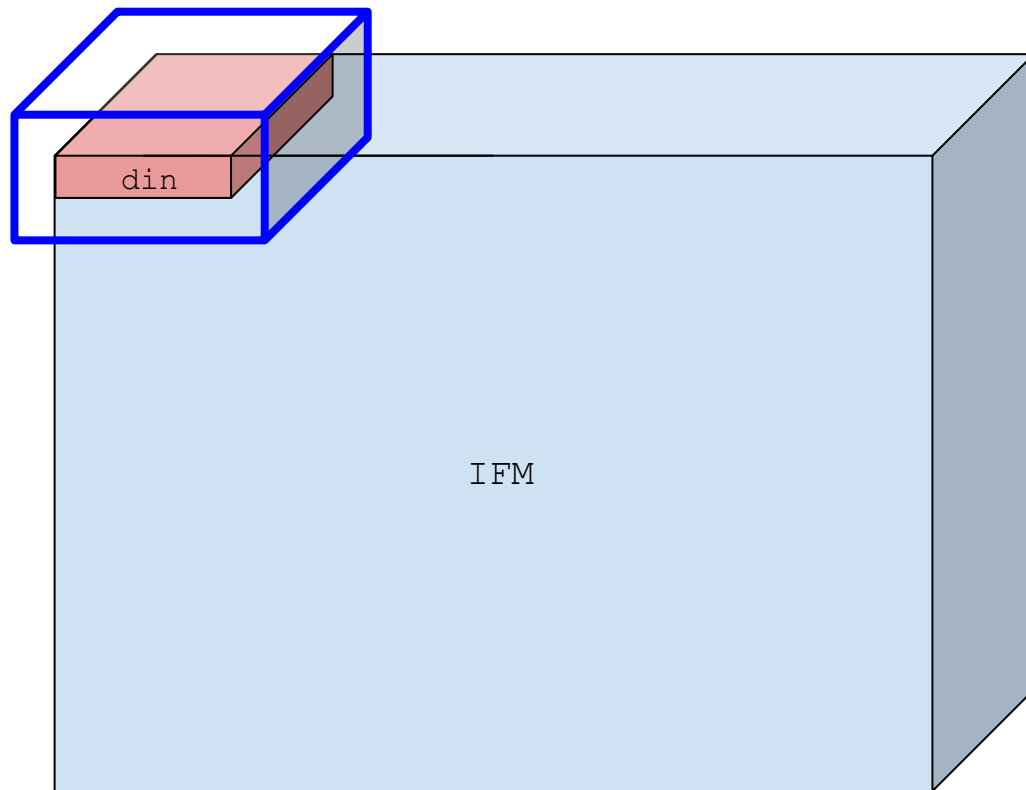
`ch_batch_idx = 0`

Layer 1 Convolution Operations

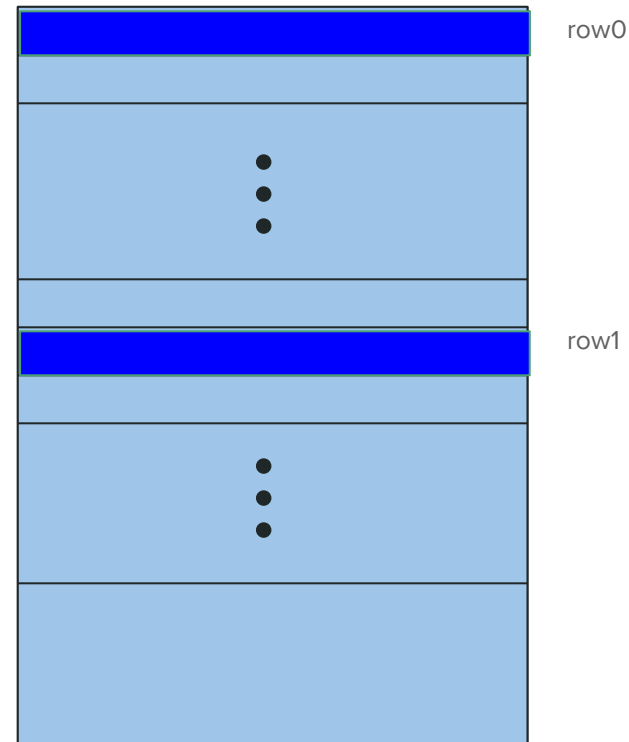


`ch_batch_idx = 1`

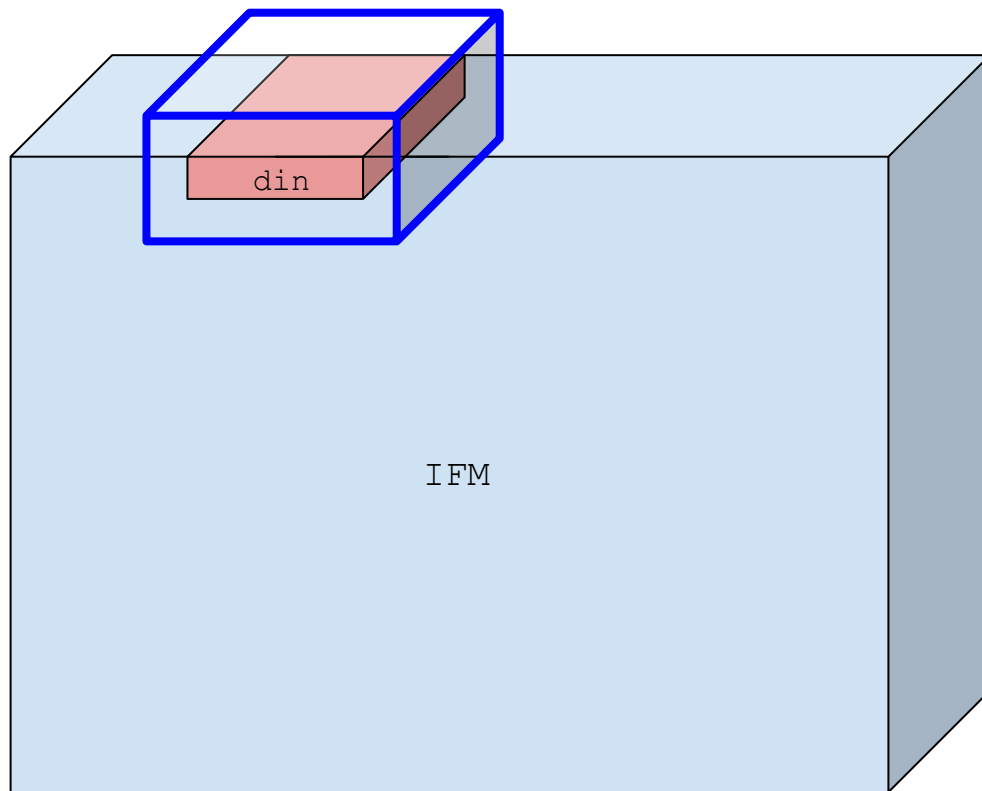
Layer 2 Convolution Operations



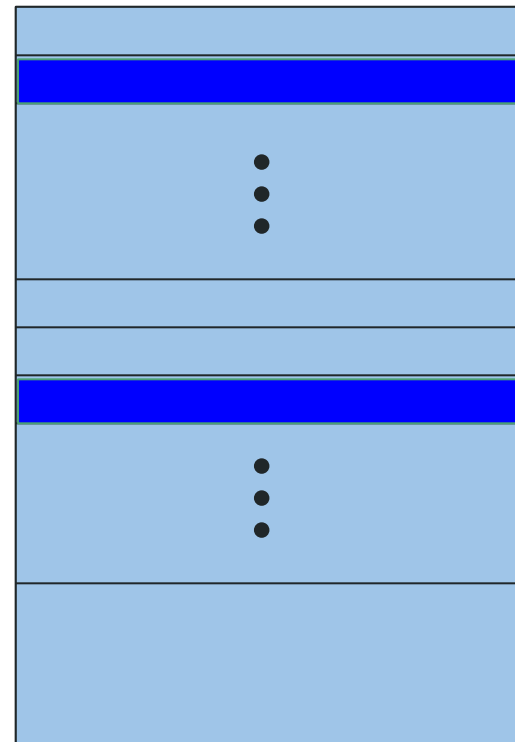
Layer 2 line buffer



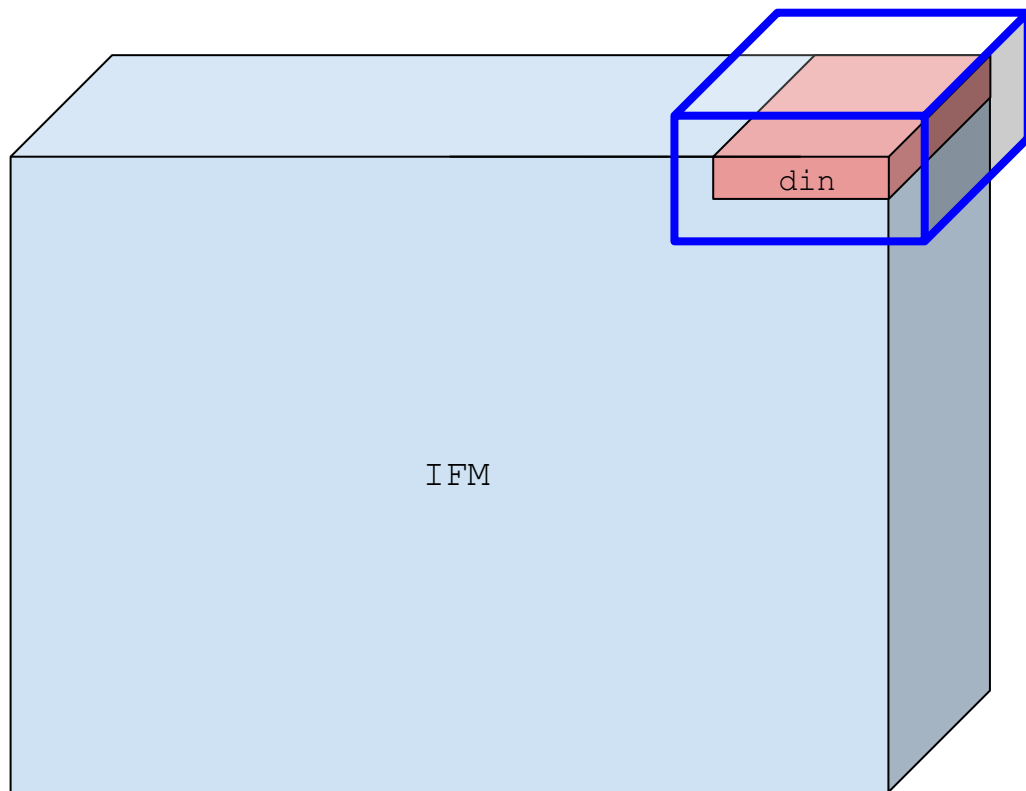
Layer 2 Convolution Operations



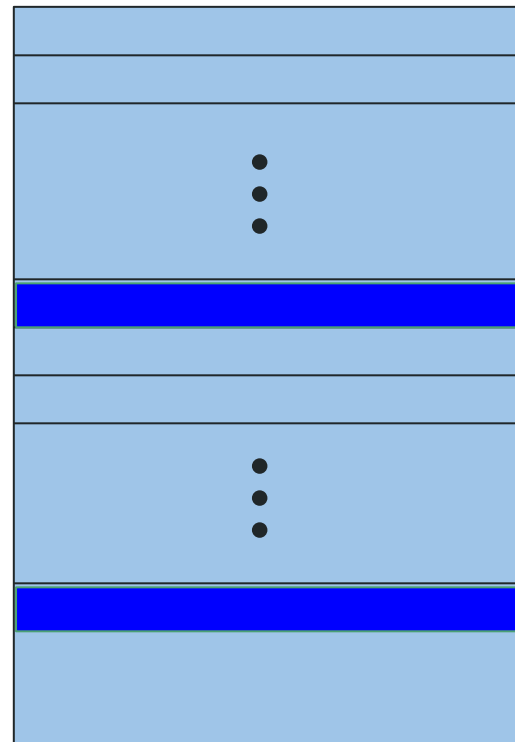
Layer 2 line buffer



Layer 2 Convolution Operations



Layer 2 line buffer



Kernel outputs structure

